

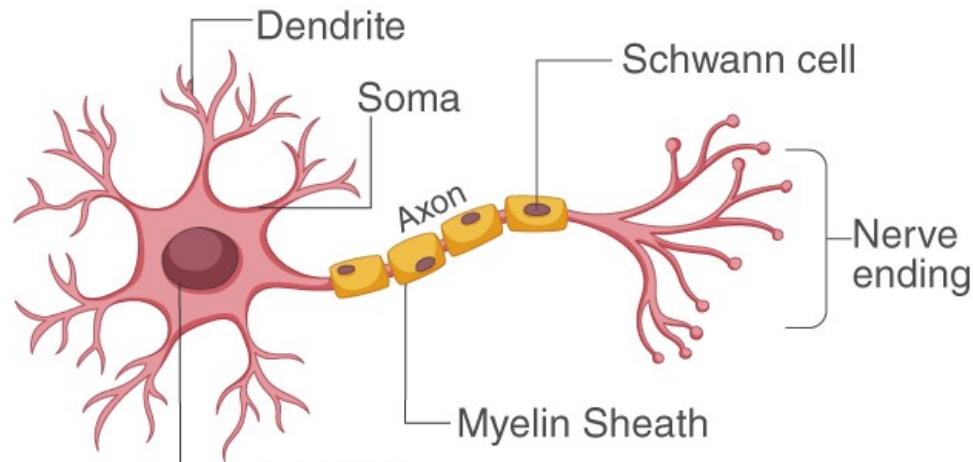


Discovering AI

Concept level



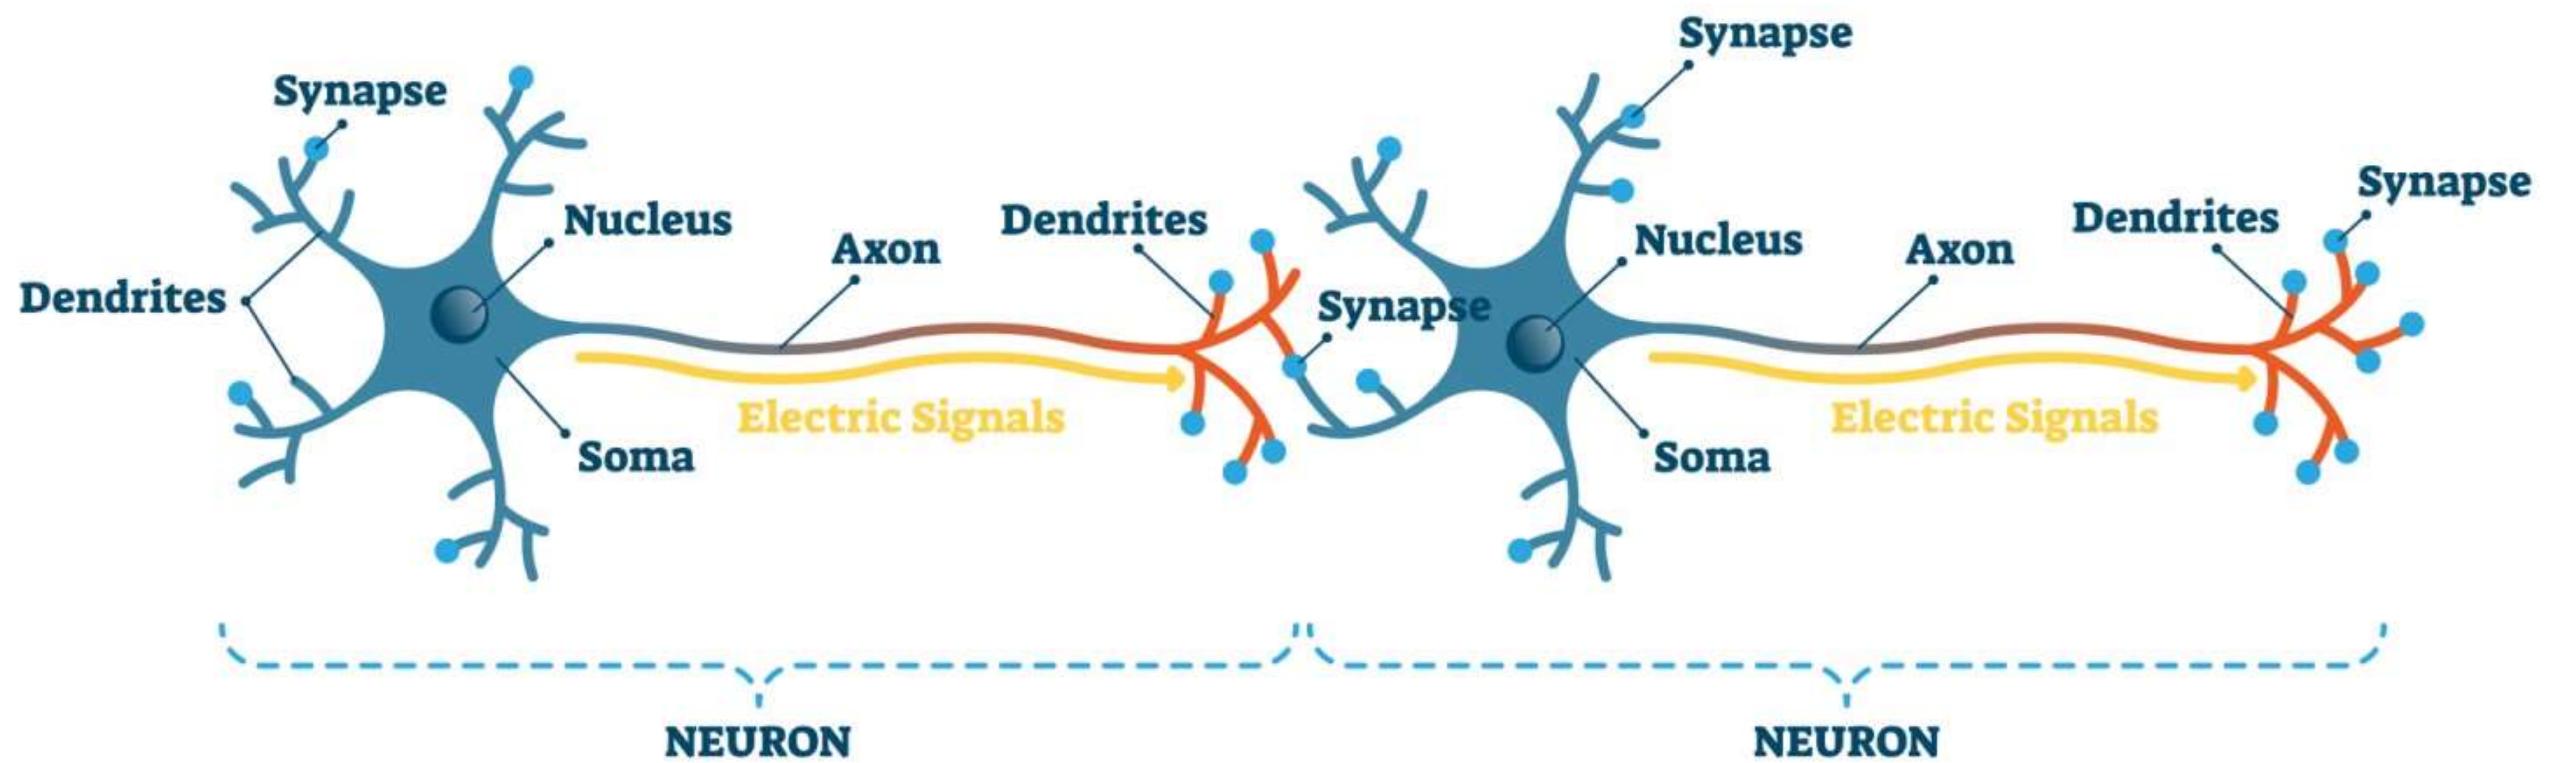
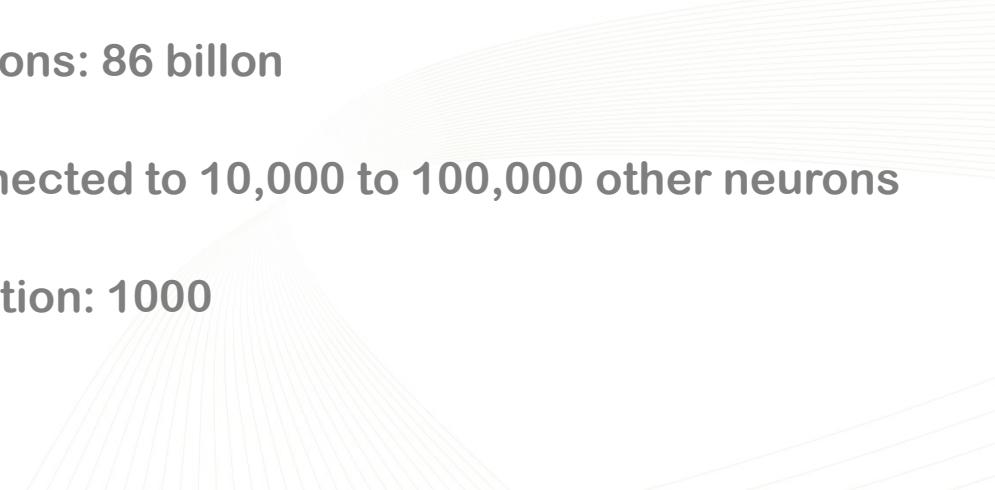
STRUCTURE OF NEURON

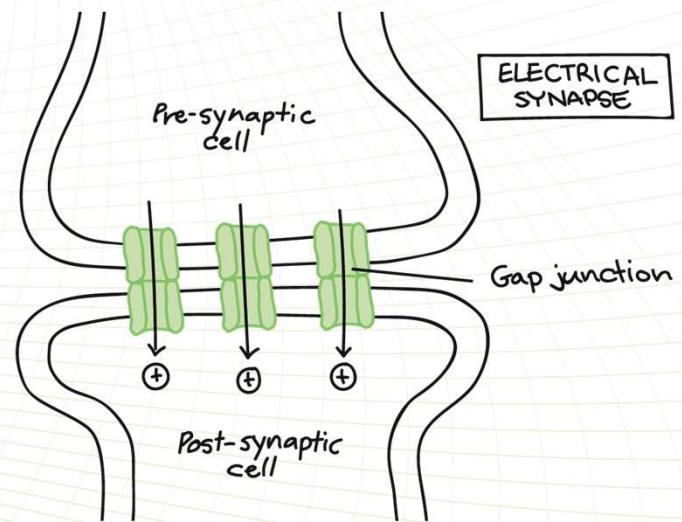
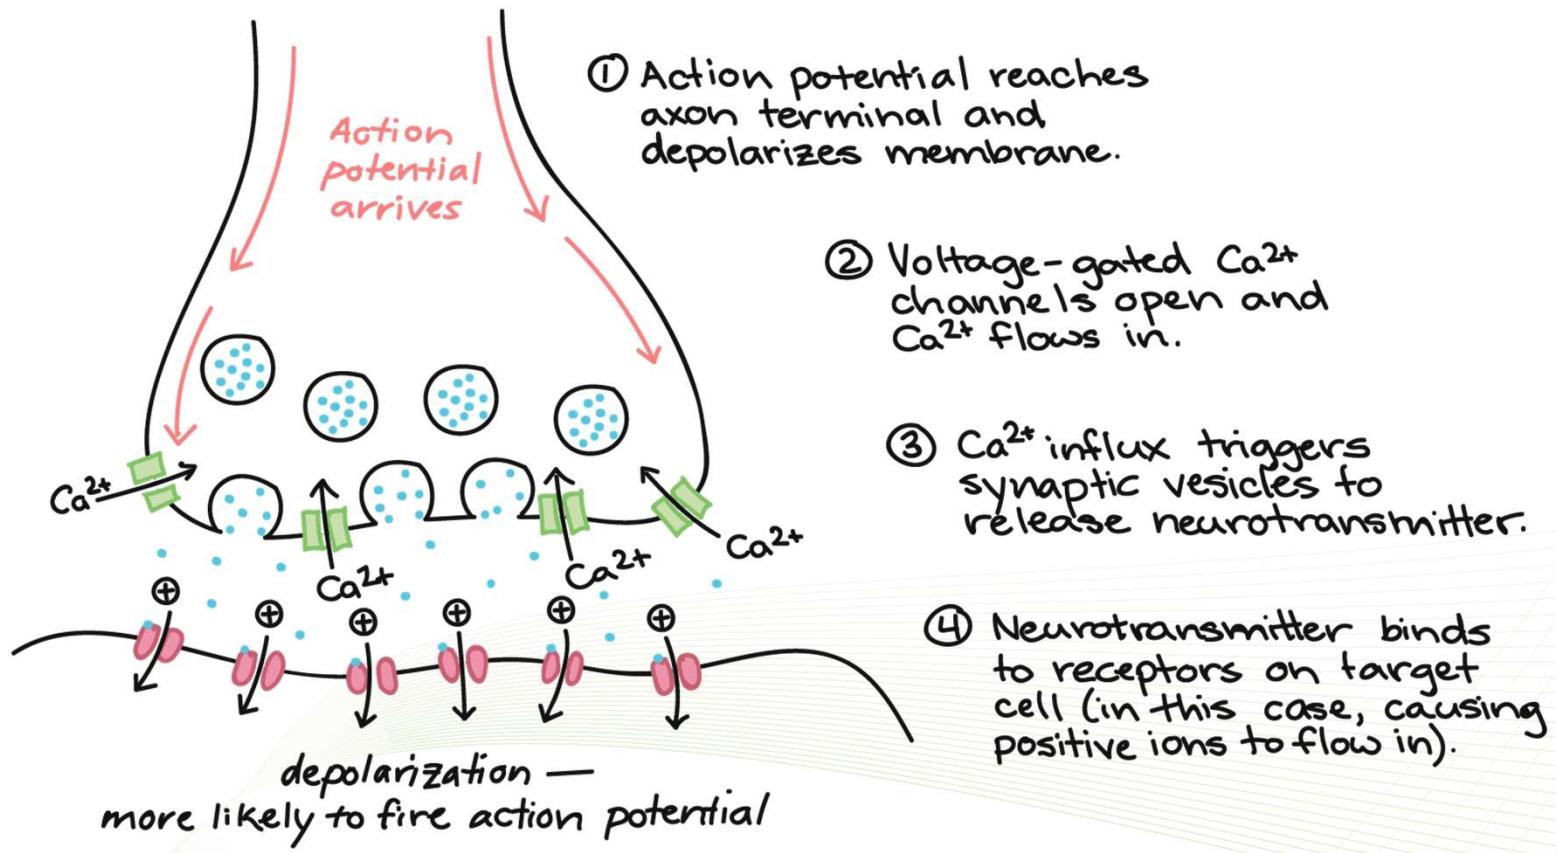


Number of Neurons: 86 billion

A neuron is connected to 10,000 to 100,000 other neurons

Average connection: 1000

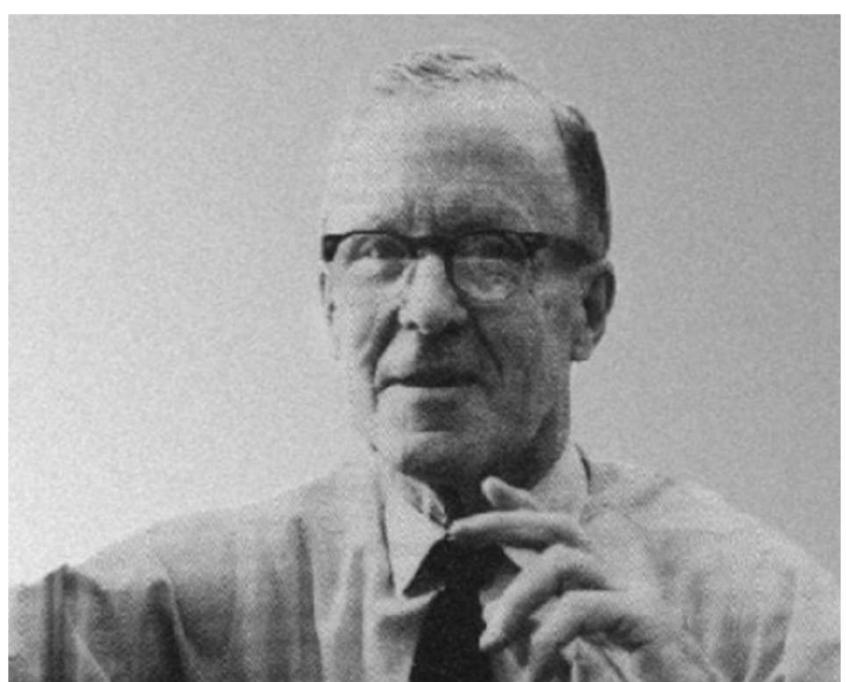
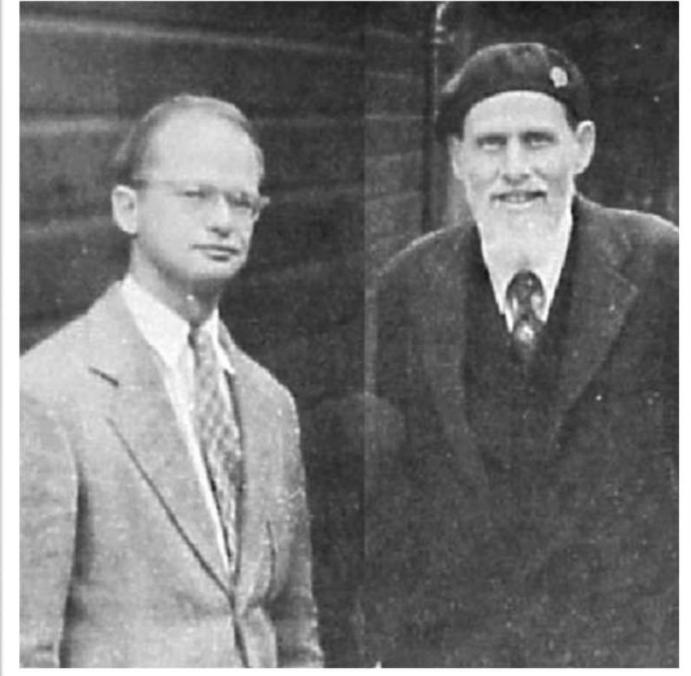




History

Walter Pitts and Warren McCulloch (1943)

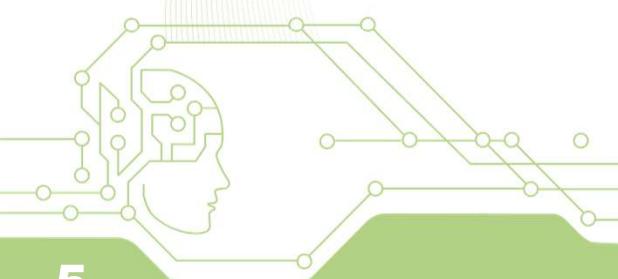
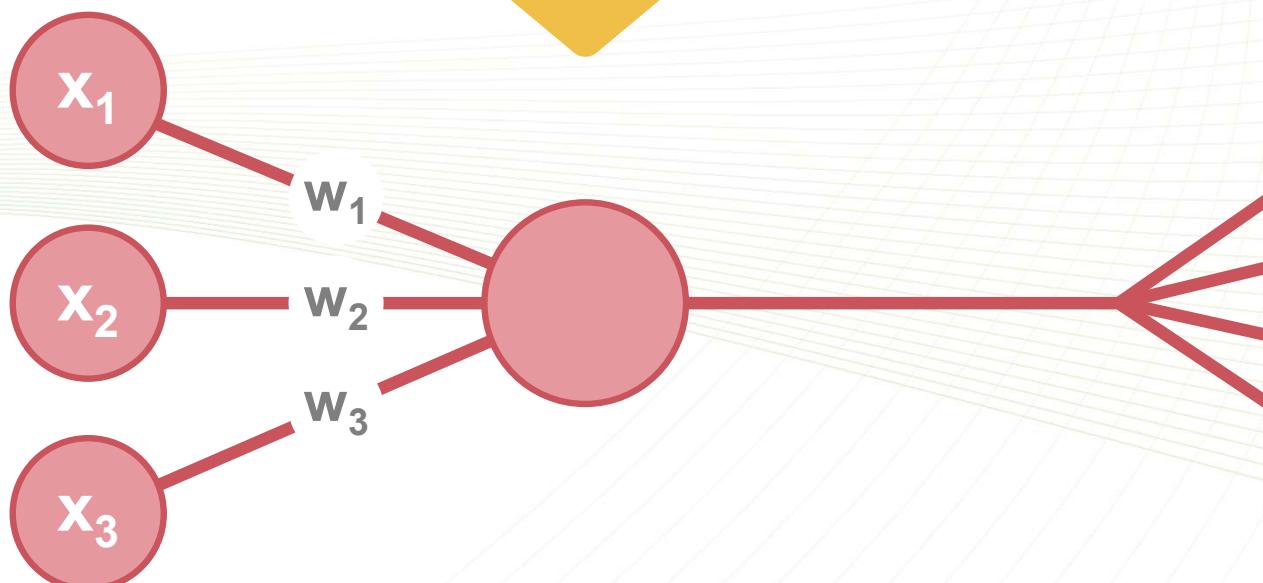
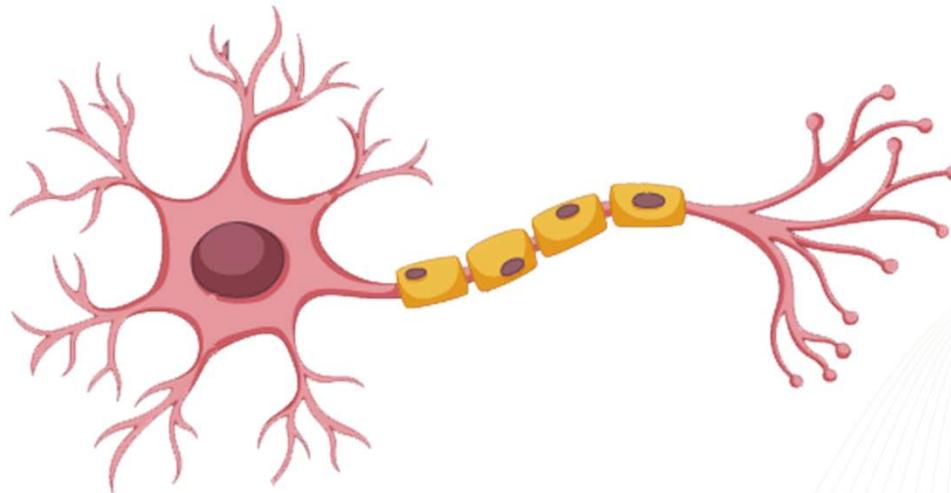
- First model of neurons with the On/Off capability
- Firing in response to sufficient stimulation from surrounding neurons



Donald Hebb (1949) Psychology

- Inventing the concept of learning in artificial neurons
- A simple method to update the connection strength between neurons

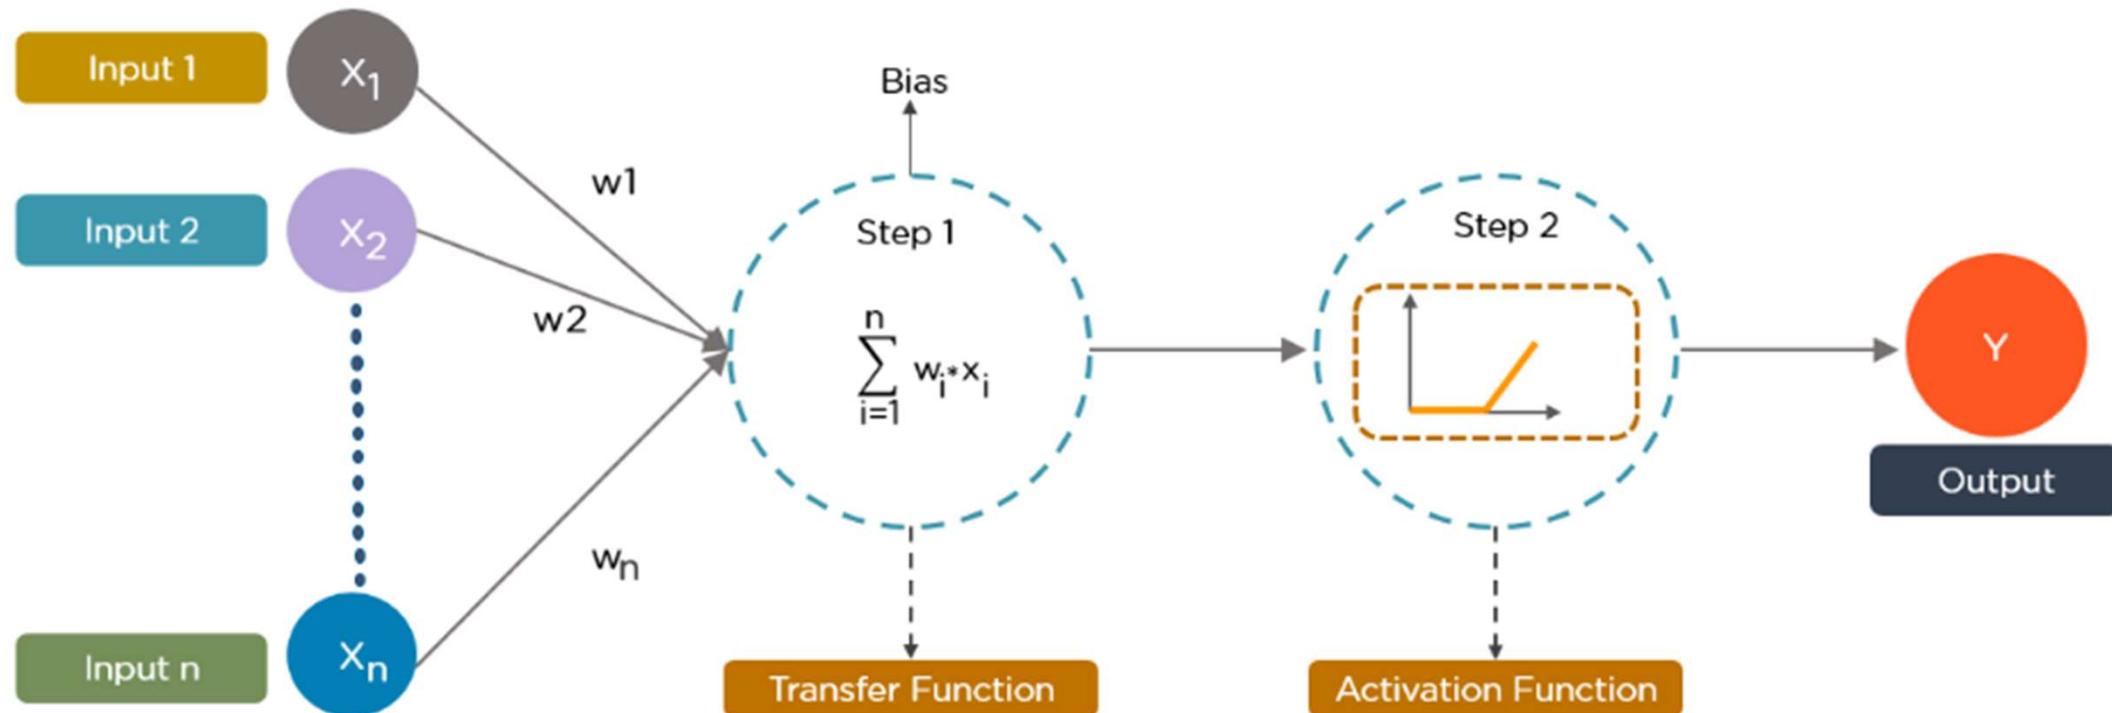
Reference: www.simplypsychology.org



5

Reference: www.simplypsychology.org

Discovering AI
Sasan Azimi



Types of Data

Tables

(Structured data) 

XML, json

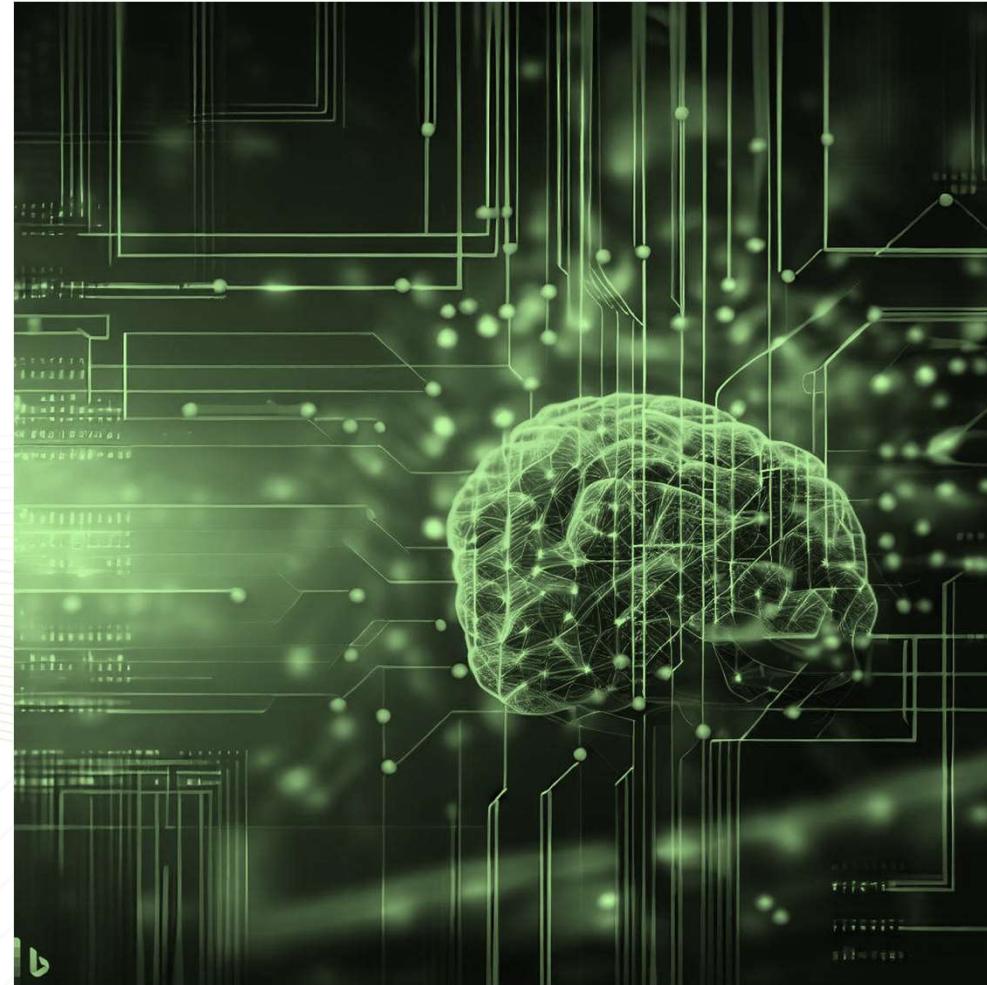
(Semi-Structured data) 

Image, audio, text

(unstructured data) 

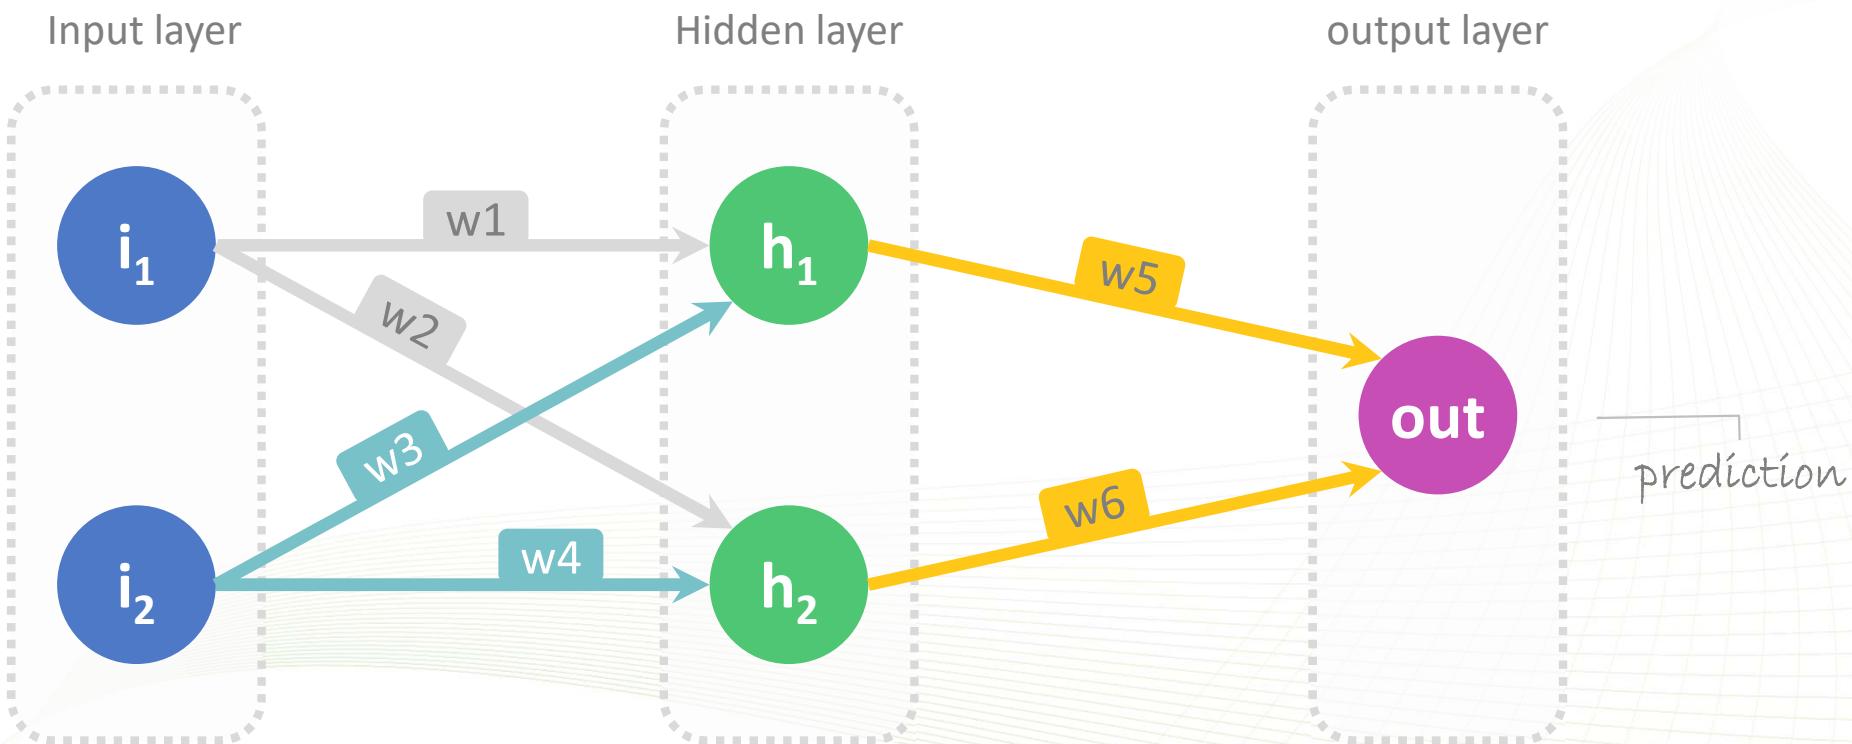
Machine learning

- Supervised Machine learning
- Unsupervised Machine learning
- Semi-supervised learning
- Reinforcement Machine learning
- Self-supervised learning





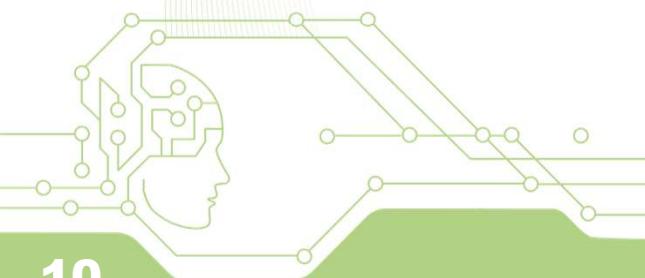
gorilla fishing in the ocean,

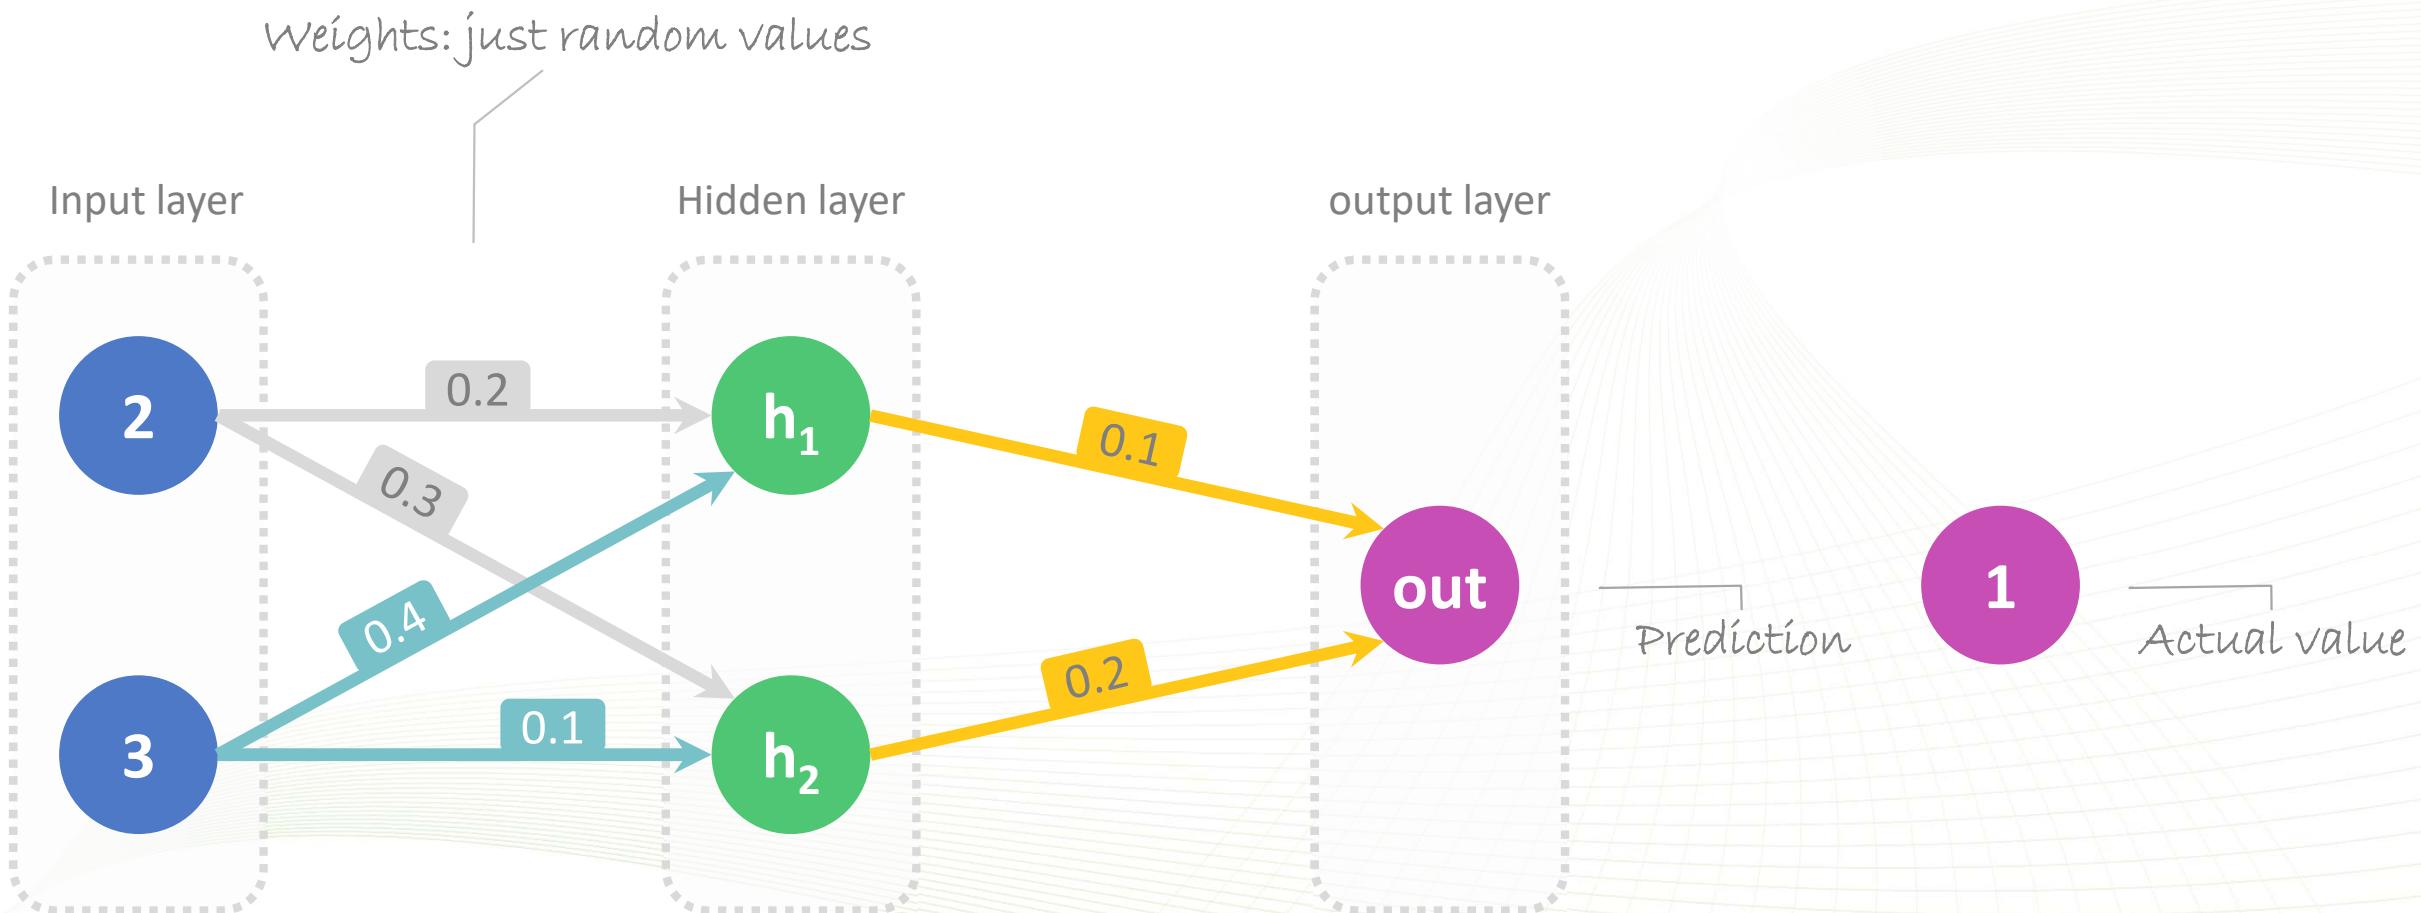


$$h_1 = i_1 w_1 + i_2 w_3$$

$$h_2 = i_1 w_2 + i_2 w_4$$

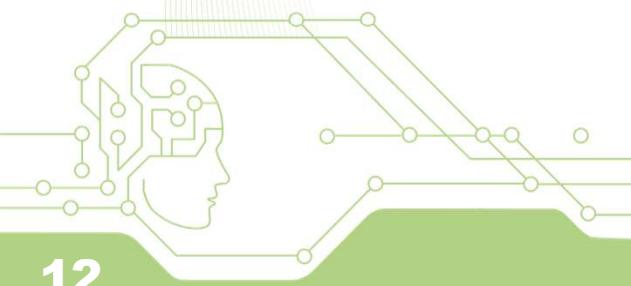
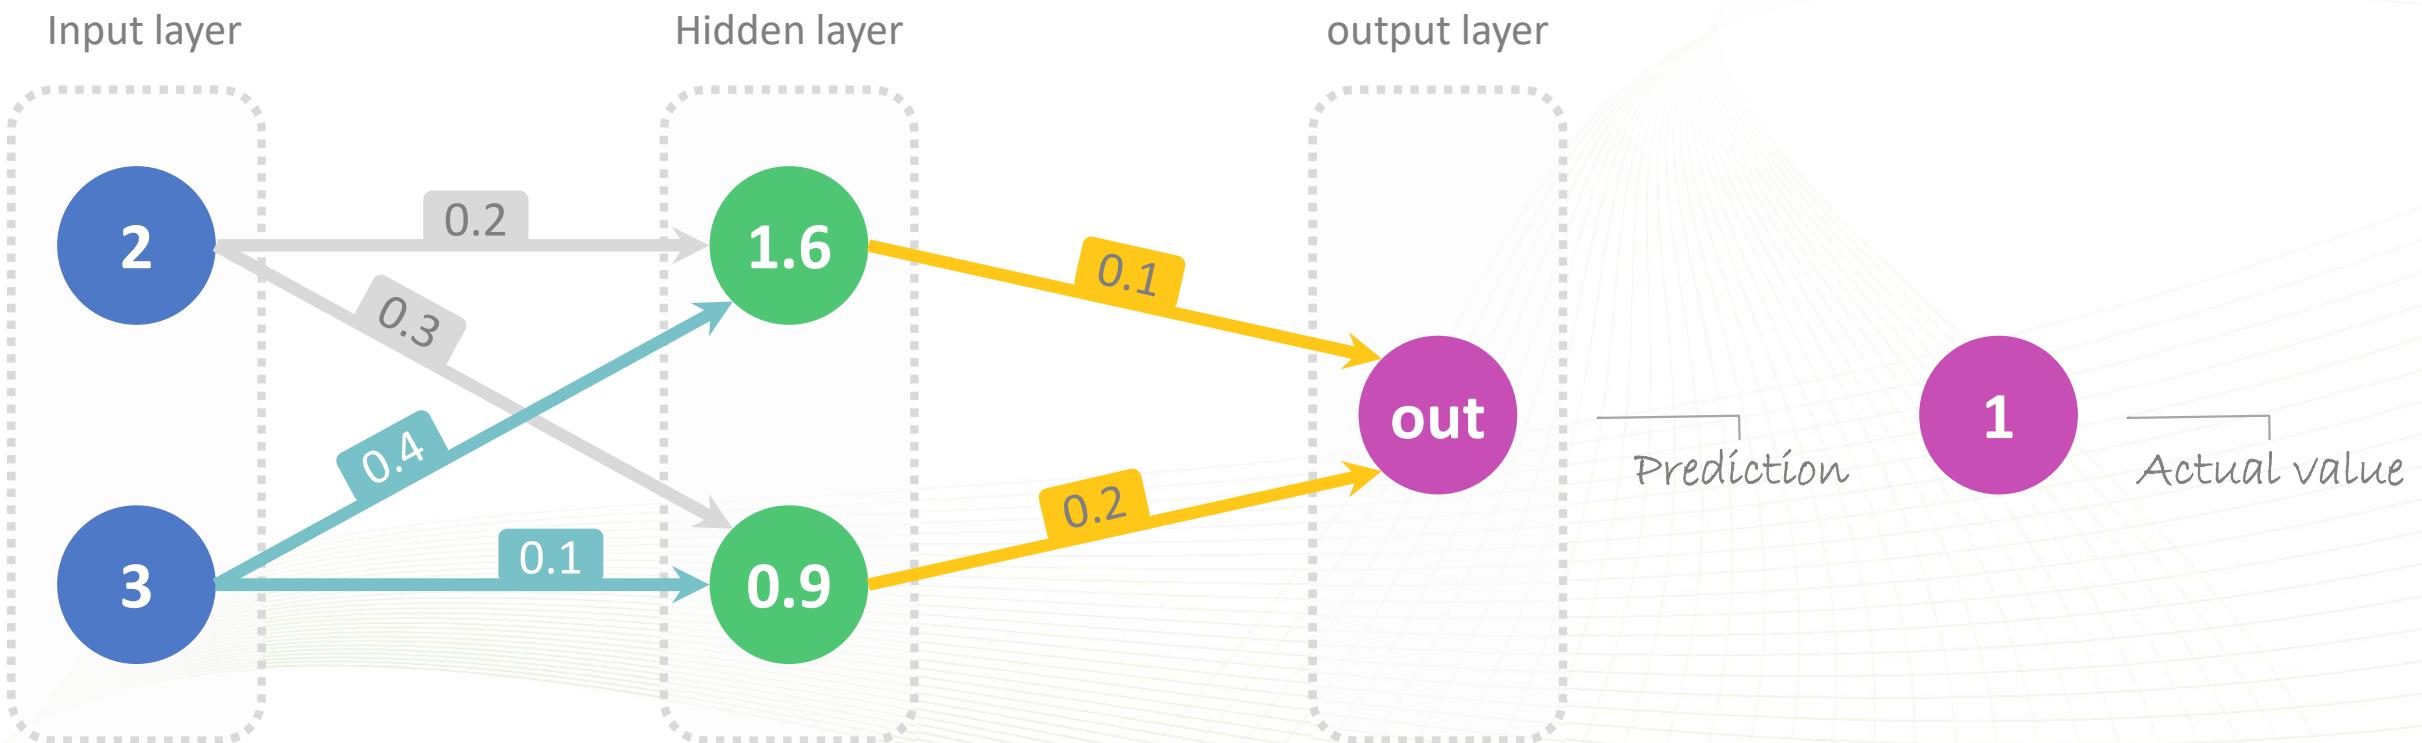
$$out = h_1 w_5 + h_2 w_6$$

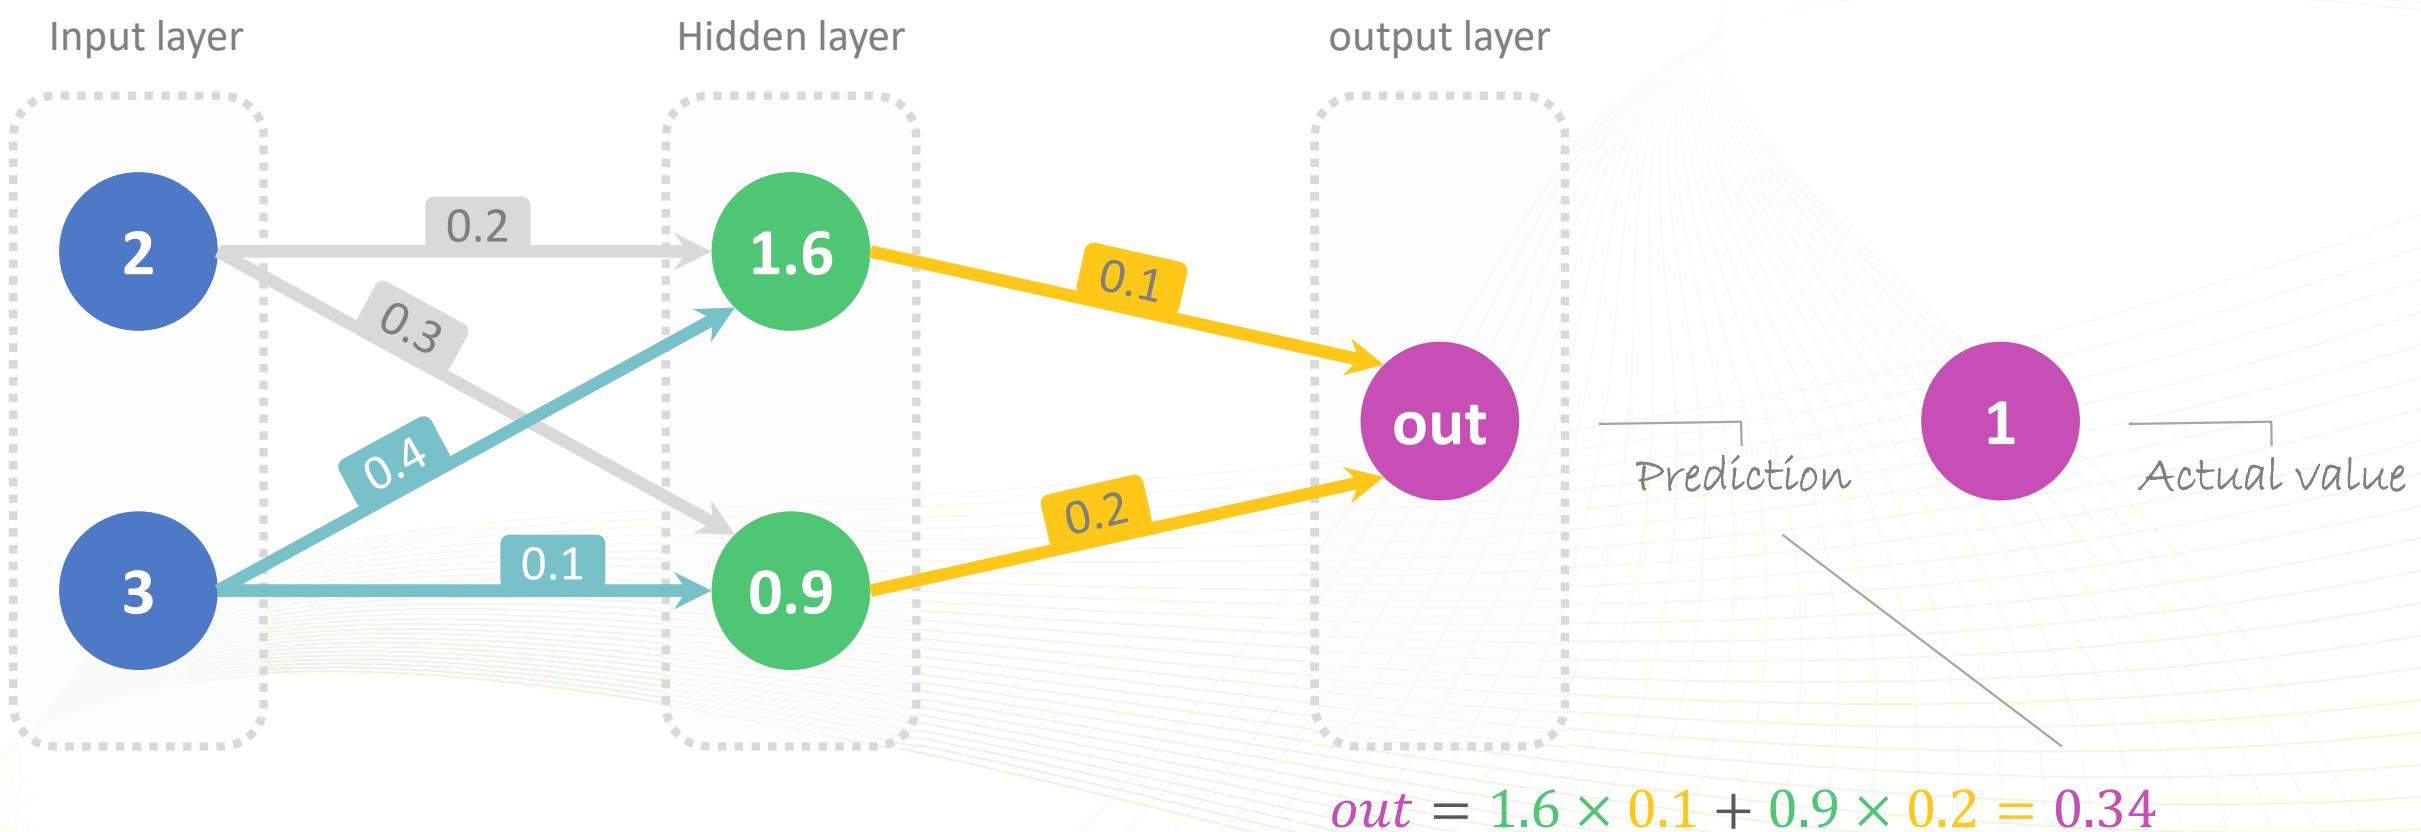




$$h_1 = 2 \times 0.2 + 3 \times 0.4 = 1.6$$

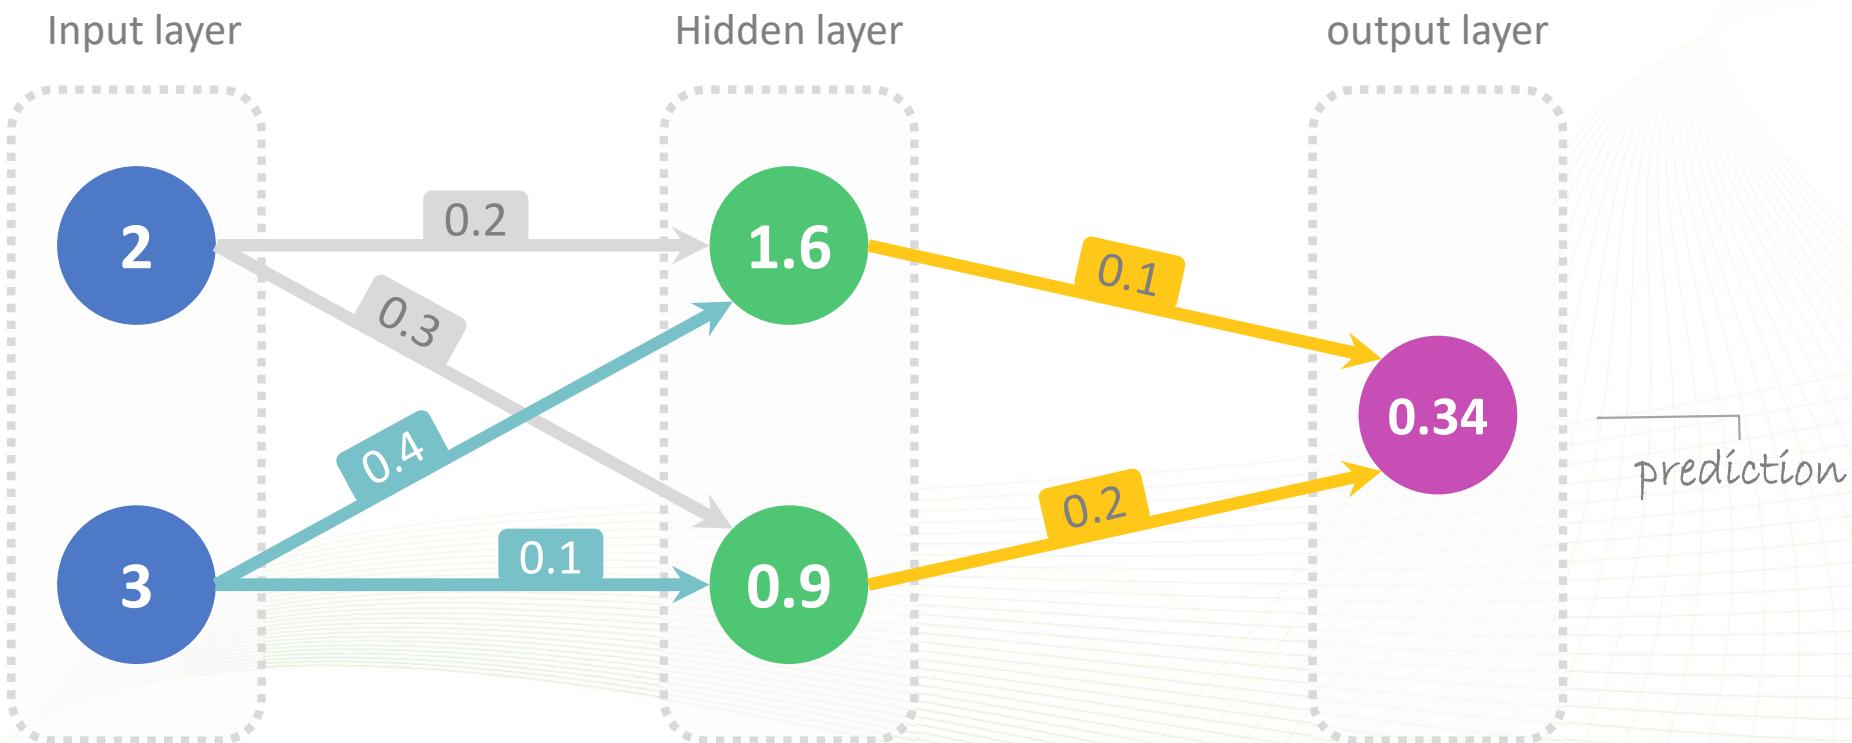
$$h_2 = 2 \times 0.3 + 3 \times 0.1 = 0.9$$





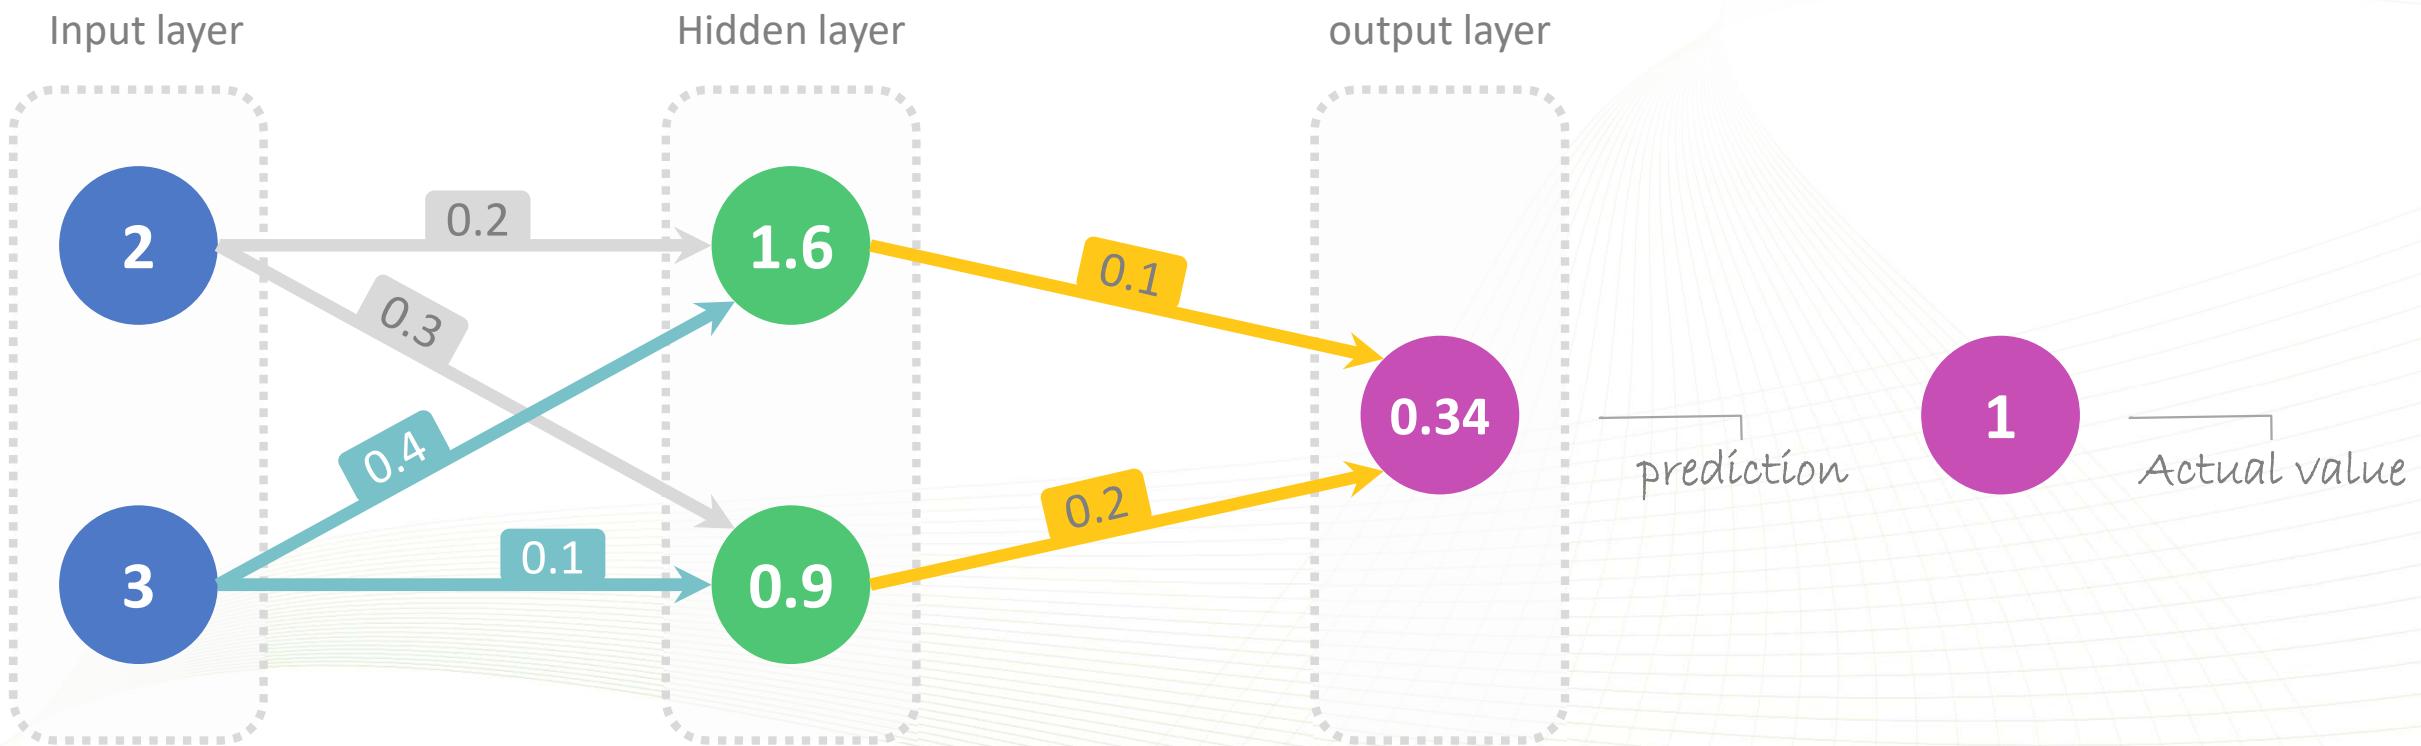
$$h_1 = 2 \times 0.2 + 3 \times 0.4 = 1.6$$

$$h_2 = 2 \times 0.3 + 3 \times 0.1 = 0.9$$



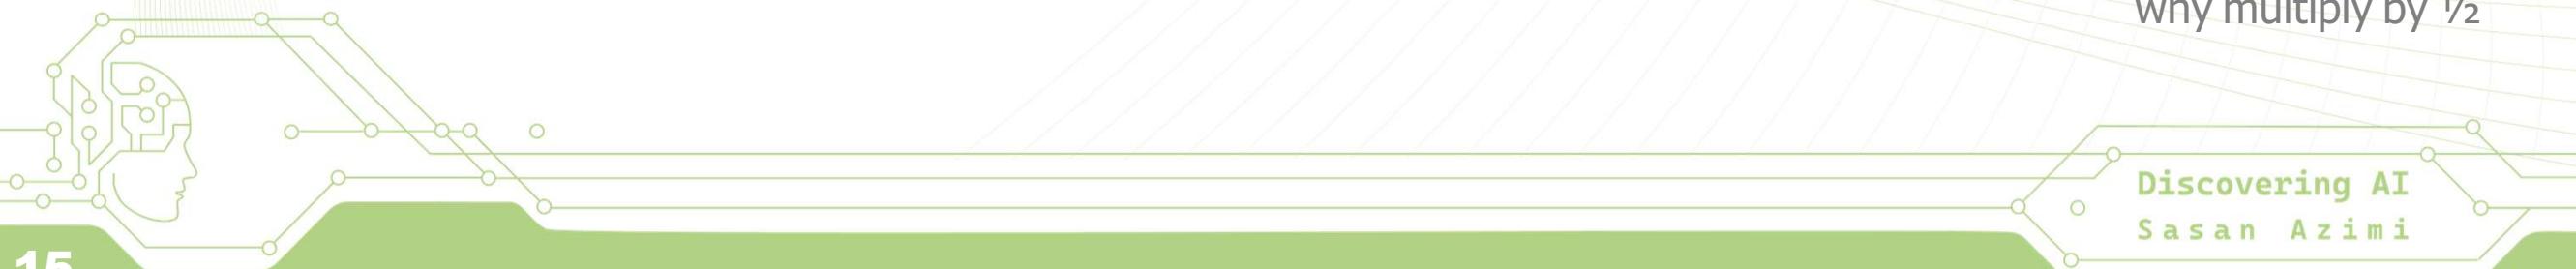
$$h_1 = 2 \times 0.2 + 3 \times 0.4 = 1.6$$

$$h_2 = 2 \times 0.3 + 3 \times 0.1 = 0.9$$



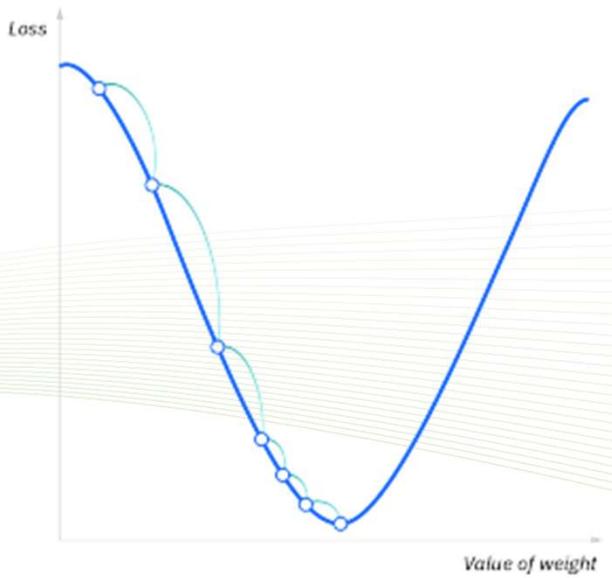
$$\text{error} = \frac{1}{2} (\text{prediction} - \text{actual value})^2$$

$$\frac{1}{2} (0.34 - 1)^2 = 0.218$$

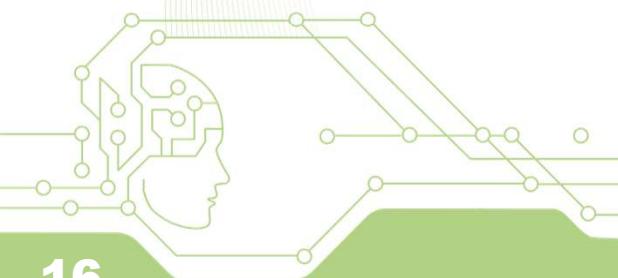
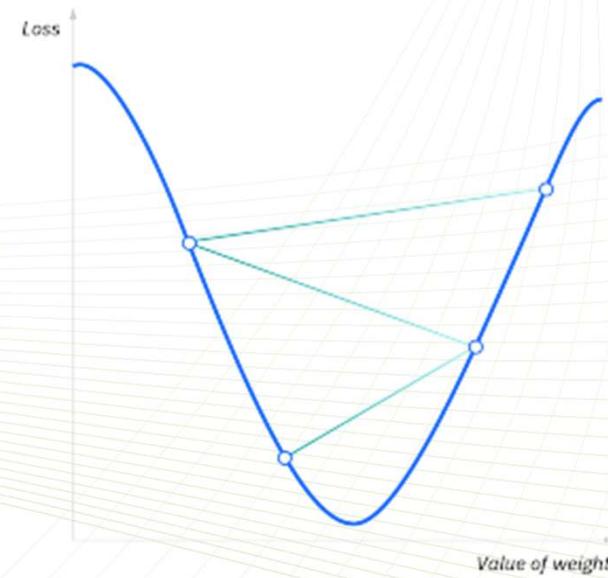


Learning rate

Small learning rate



Large learning rate



Backpropagation

Target: minimizing the error

How could we change the prediction?

Using derivation to recognize the direction of movement



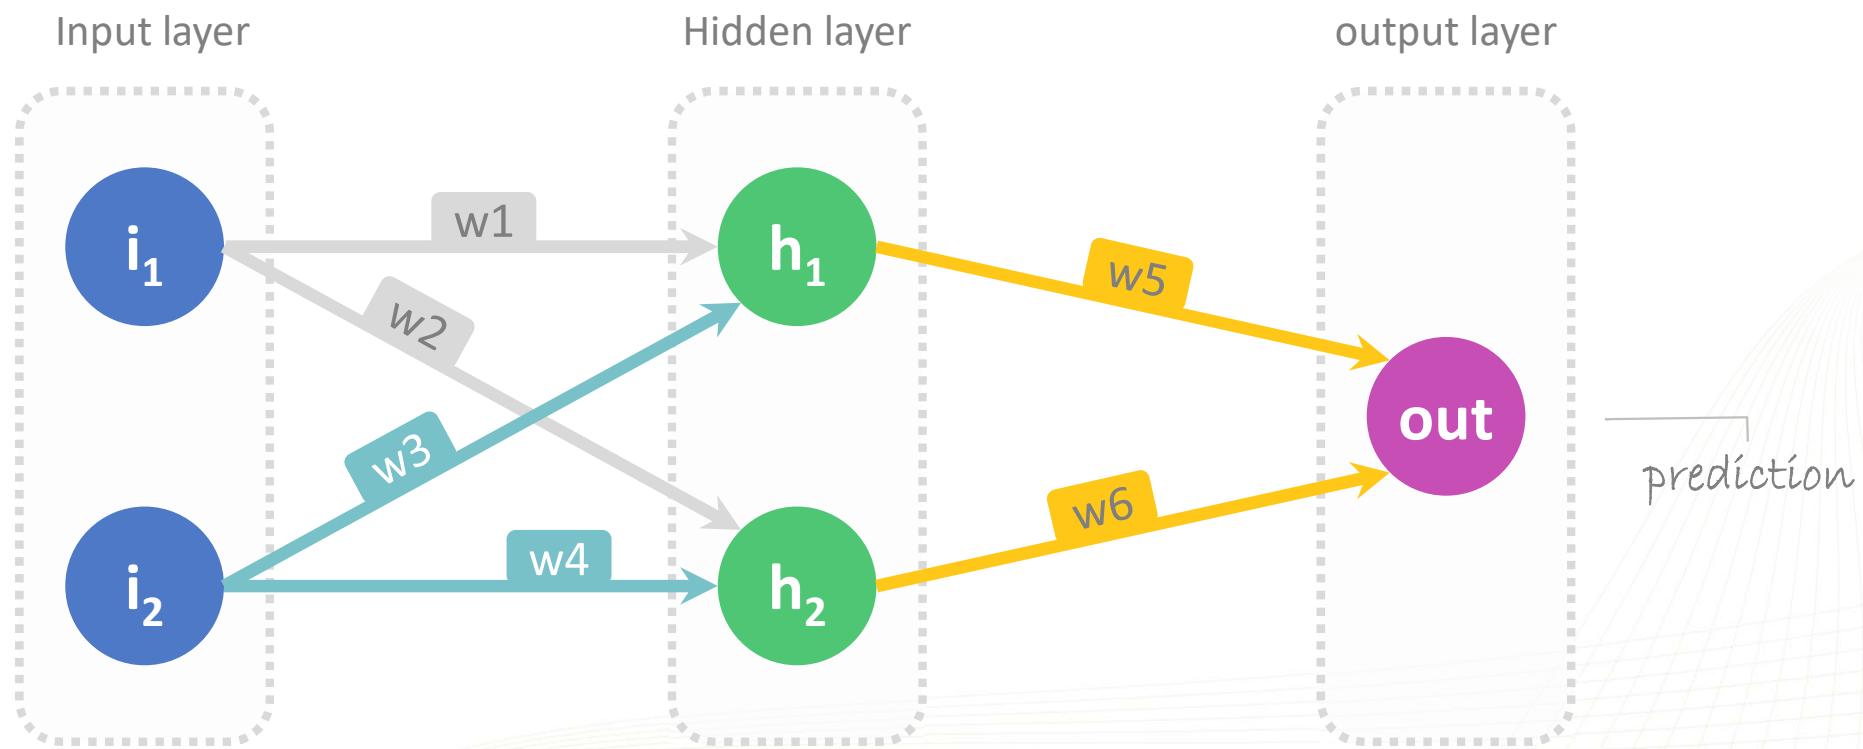
Gradient descent

Improving the prediction by changing the weights

$$w_x = w_x - \alpha \left(\frac{\partial \text{error}}{\partial w_x} \right)$$

Learning rate

$$\frac{\partial (x^2y^3 + xy^2)}{\partial x} = 2xy^3 + y^2$$



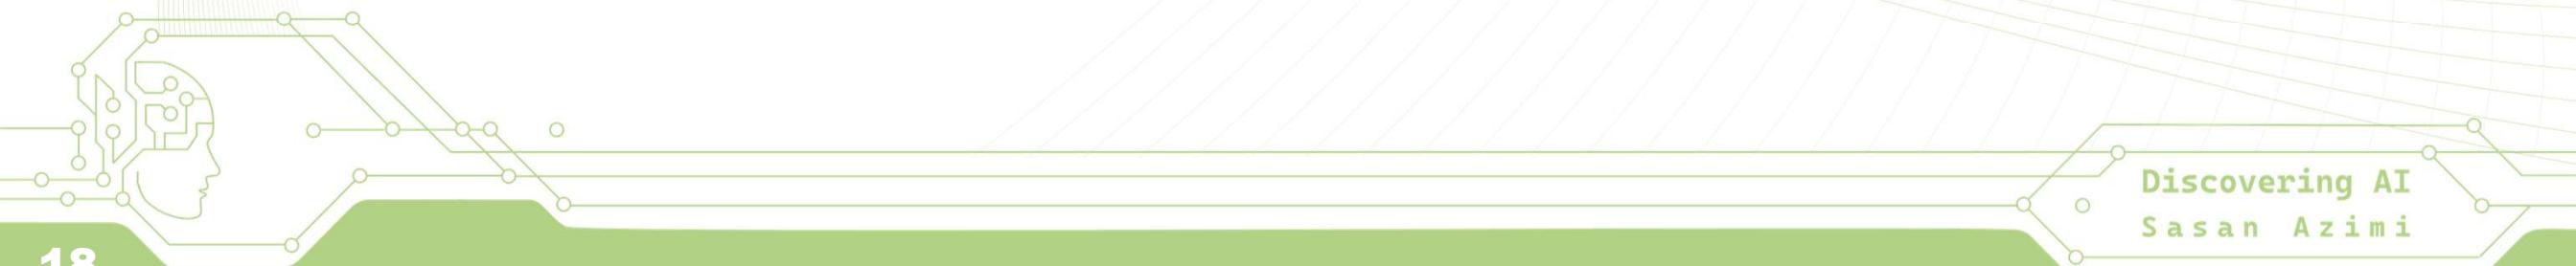
$$h_1 = i_1 w_1 + i_2 w_3$$

$$h_2 = i_1 w_2 + i_2 w_4$$

$$out = h_1 w_5 + h_2 w_6$$

$$out = (i_1 w_1 + i_2 w_3) w_5 + (i_1 w_2 + i_2 w_4) w_6$$

i_1 & i_2 fixed



$$y = u^n \longrightarrow \dot{y} = n \cdot \dot{u} \cdot u^{n-1}$$

$$w_x = w_x - \alpha \left(\frac{\partial \text{error}}{\partial w_x} \right)$$

$$w_5 = w_5 - \alpha \left(\frac{\partial \text{error}}{\partial w_5} \right)$$

$$w_5 = w_5 - \alpha(h_1 \delta)$$

$$w_6 = w_6 - \alpha(h_2 \delta)$$

$$w_4 = w_4 - \alpha(i_2 \cdot \delta w_6)$$

$$w_3 = w_3 - \alpha(i_1 \cdot \delta w_6)$$

$$w_2 = w_2 - \alpha(i_2 \cdot \delta w_5)$$

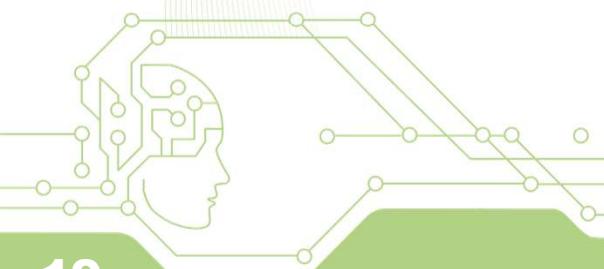
$$w_1 = w_1 - \alpha(i_1 \cdot \delta w_5)$$

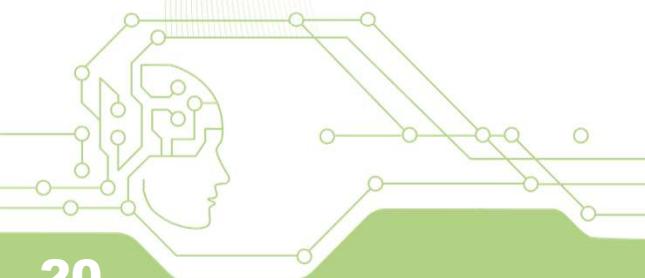
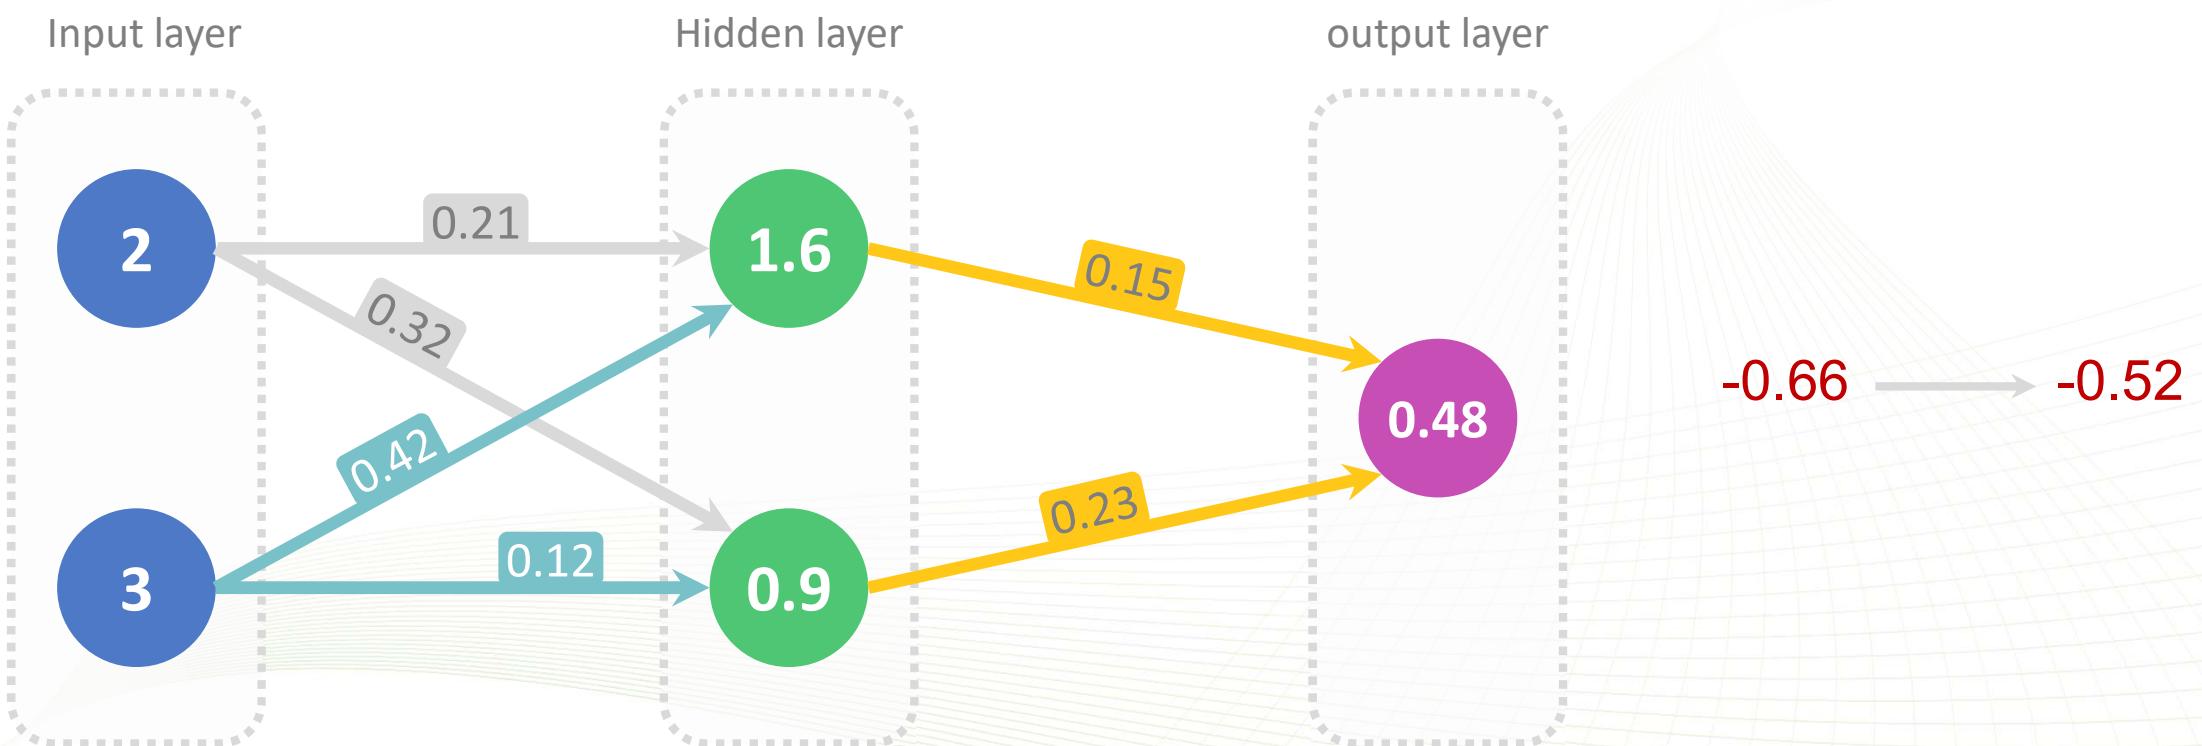
$$\text{error} = \frac{1}{2} (prediction - actual\ value)^2$$

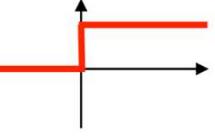
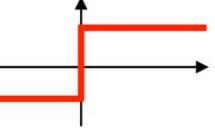
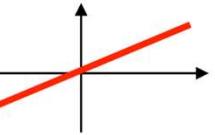
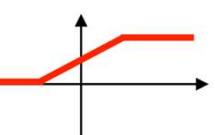
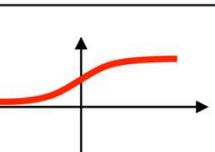
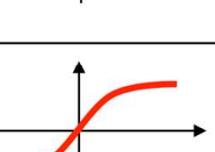
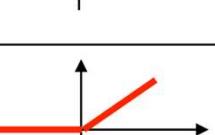
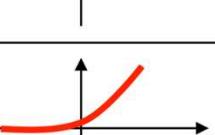
$$u = h_1 w_5 + h_2 w_6 - 1 = \delta$$

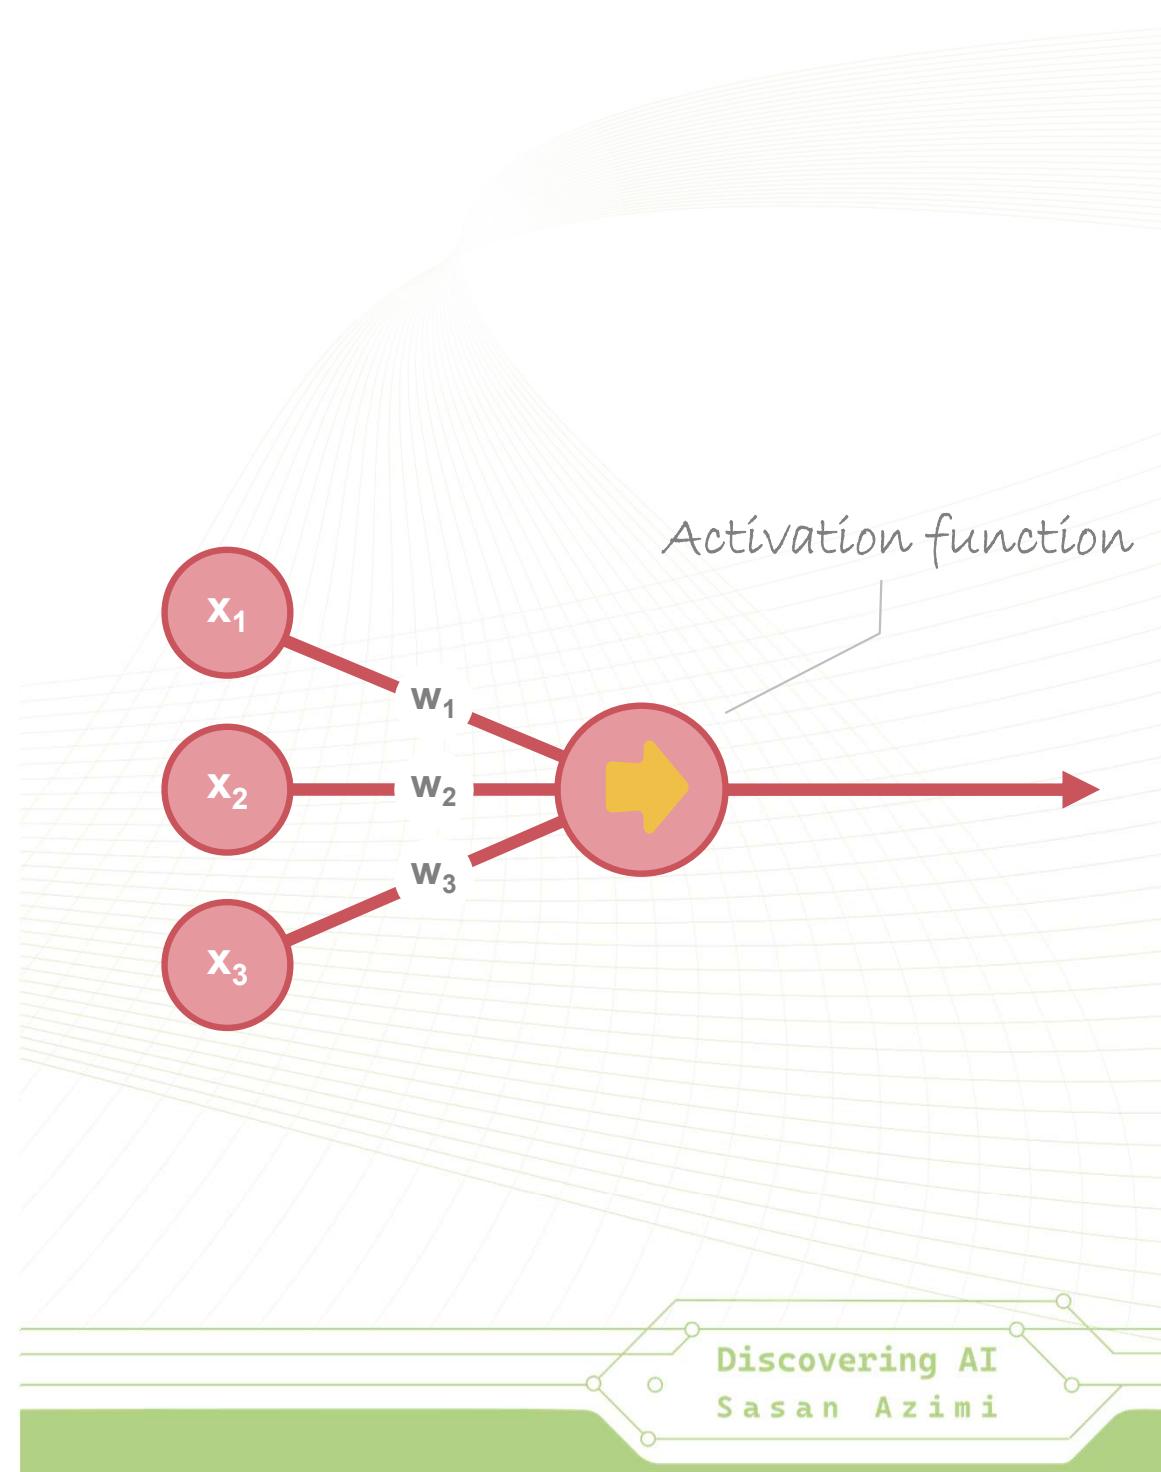
$$\text{error} = n \cdot \dot{u} \cdot u^{n-1} = \frac{1}{2} 2h_1 \delta$$

$$= h_1 \delta$$

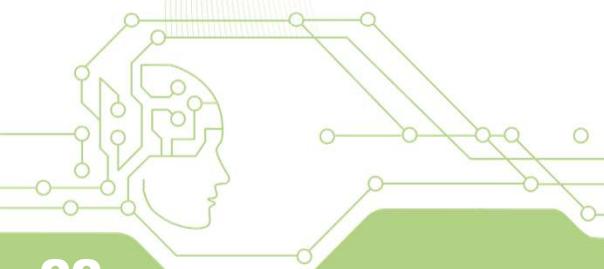
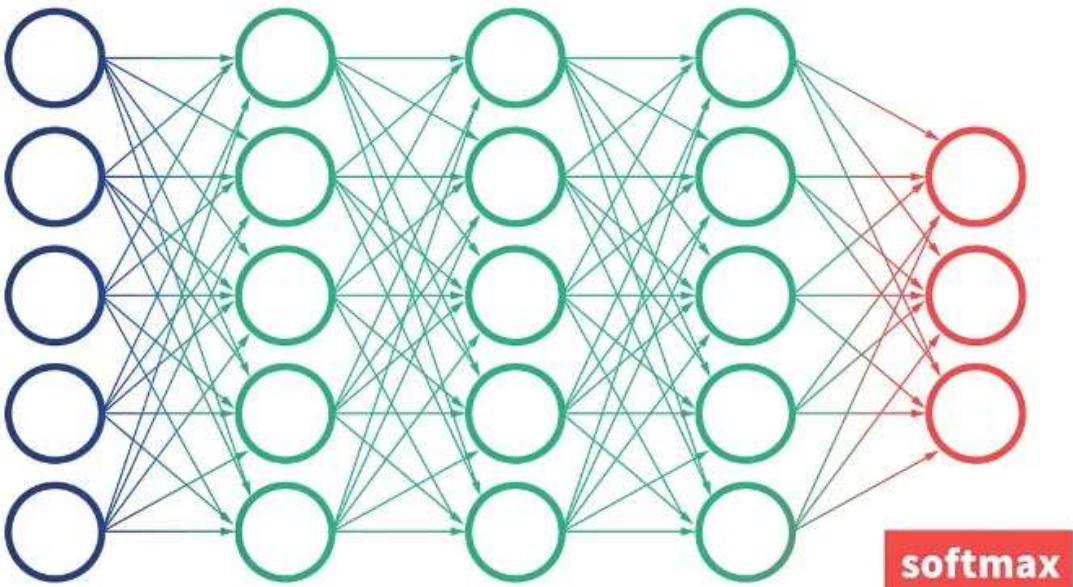




Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	



Softmax Function



Numerical conversions with softmax functions

Input	Output
5	0.730 (73.0%)
4	0.268 (26.8%)
-1	0.002 (0.2%)

The softmax function converts an input value to an output value of "0-1 values, summing to 1".
This property is used in the output layer of deep learning.

- **Softmax Formula**

$$s(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

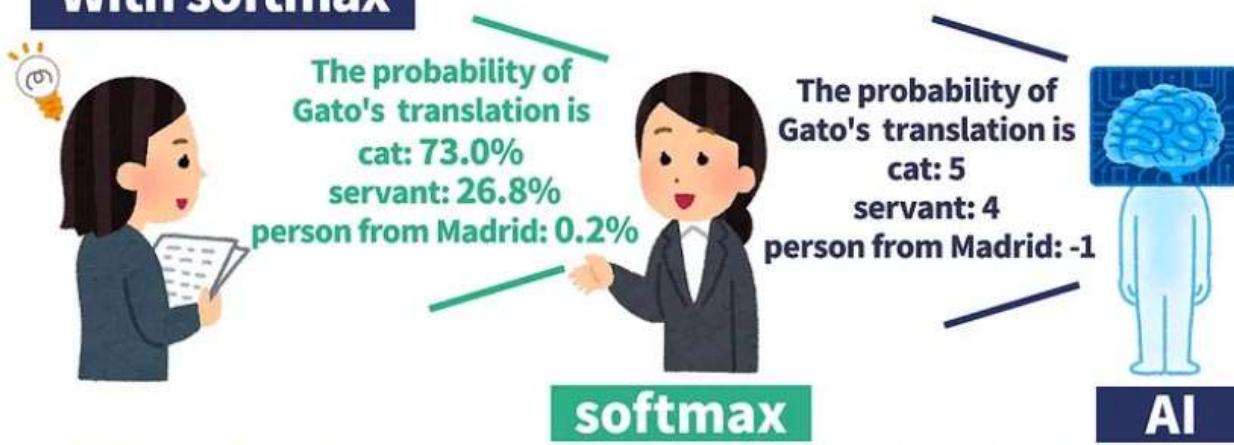


Reasons to use softmax function (Task to translate 'gato')

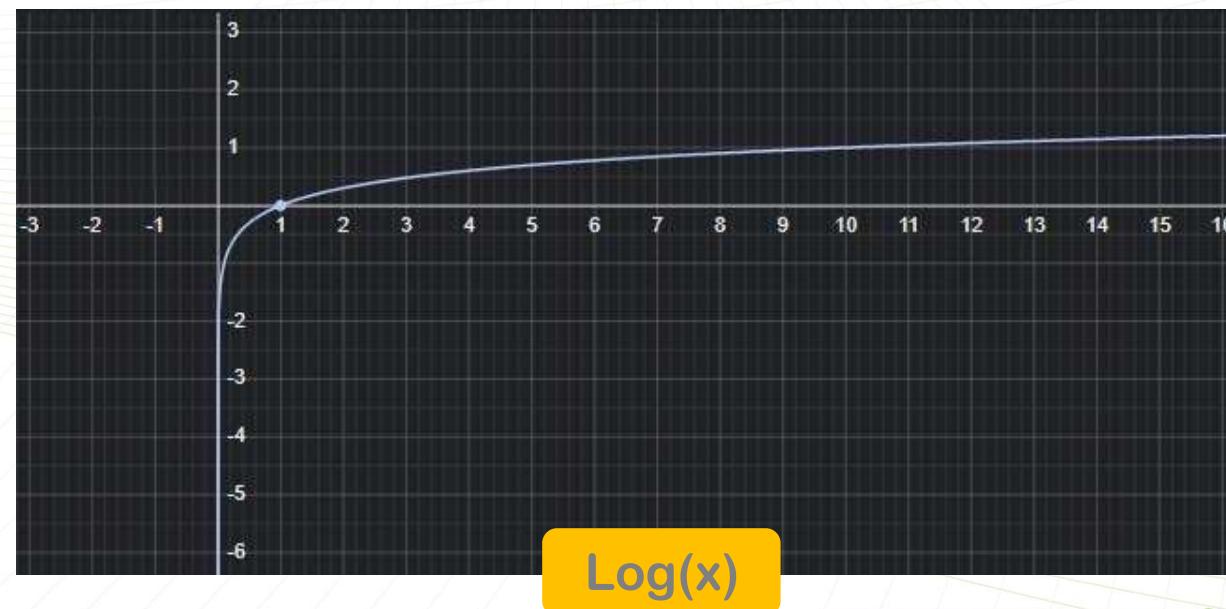
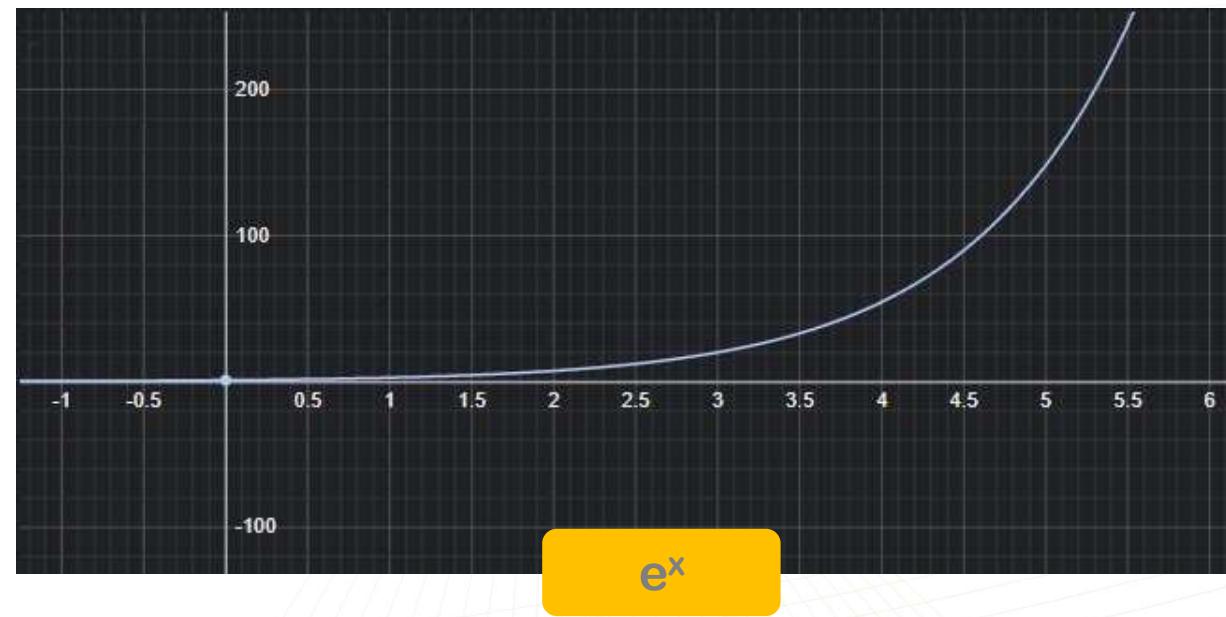
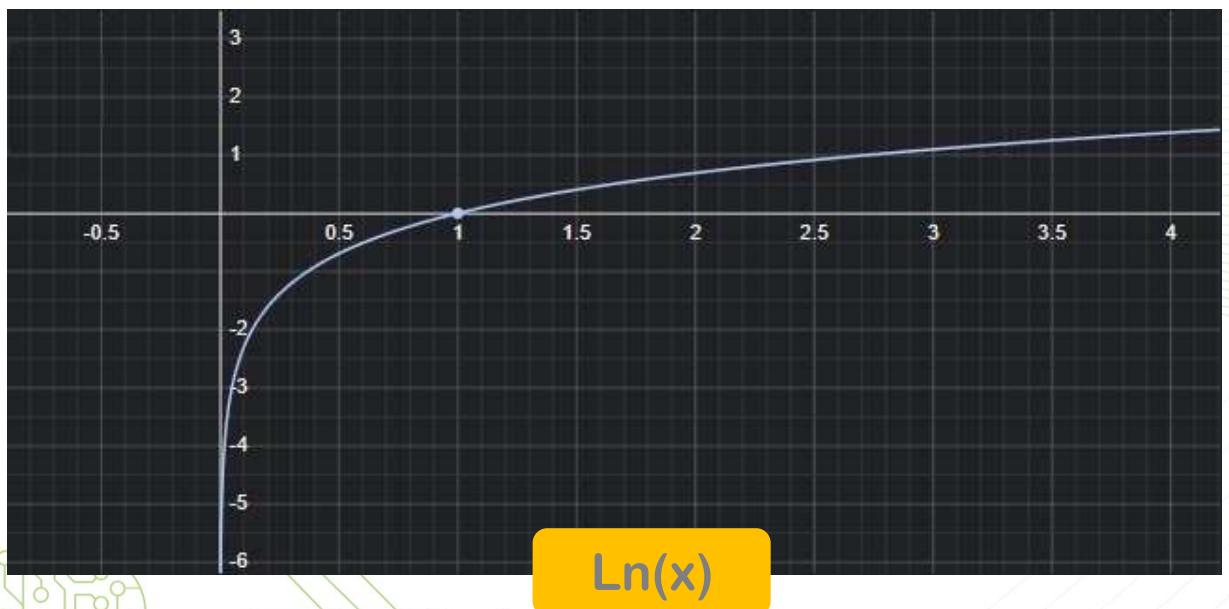
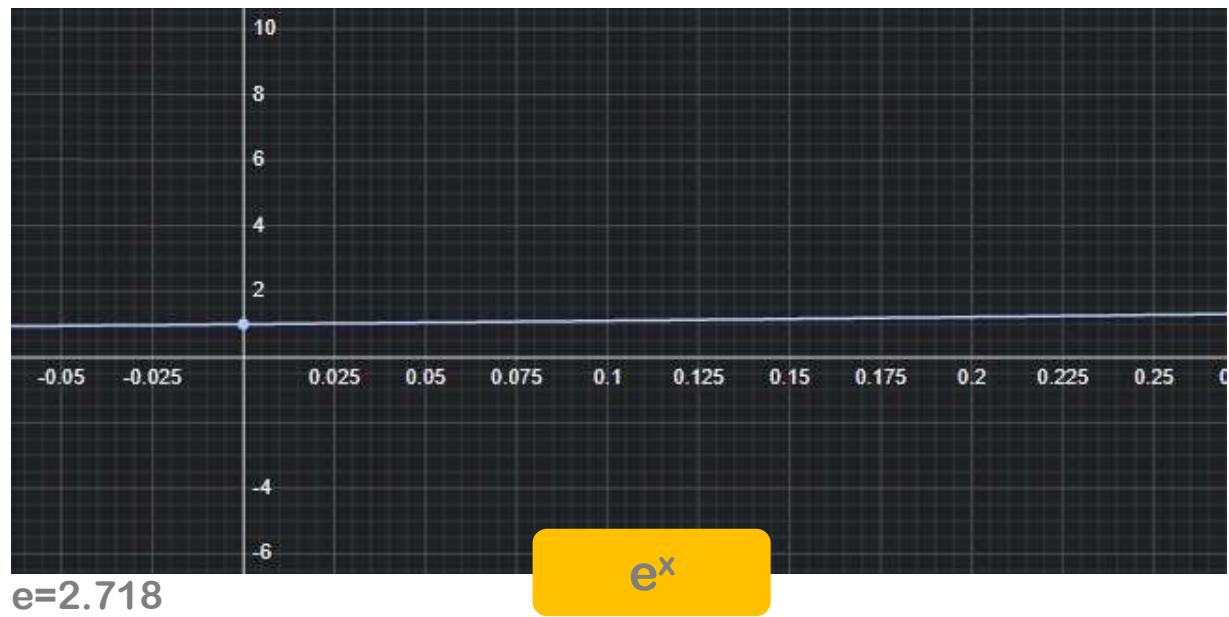
Without softmax



With softmax



Softmax function can convert to intuitively understandable probability values.

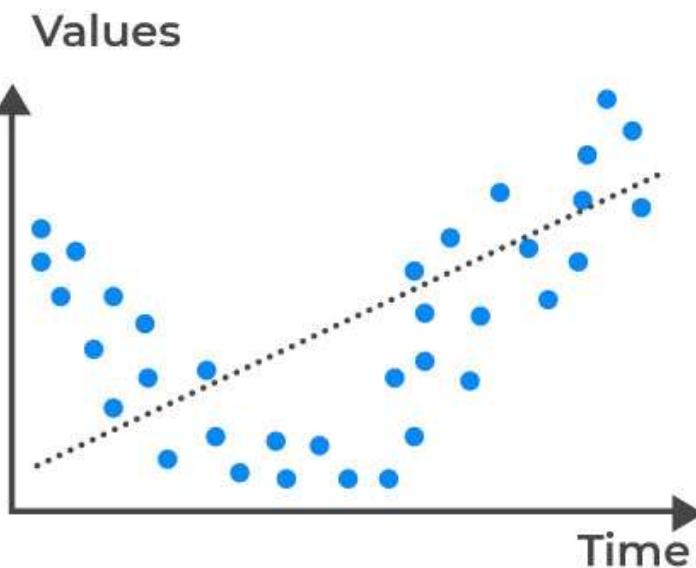


Overfit & underfit

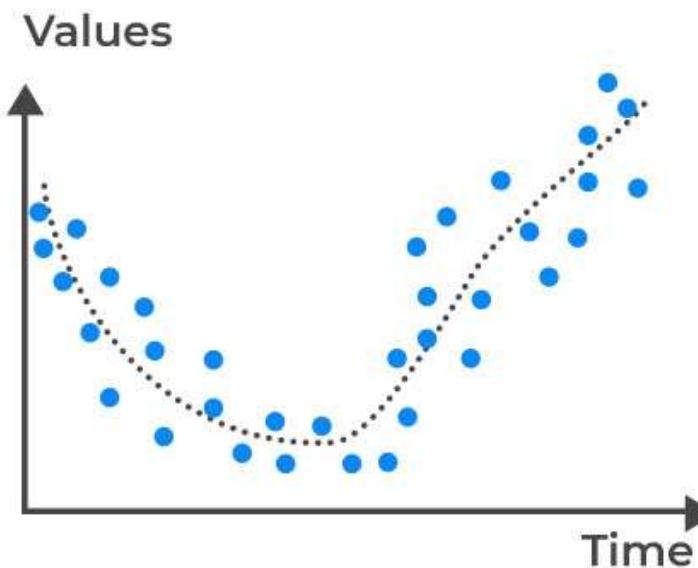
overfitting (high variance):

when the model fits perfectly to the training examples => limited generalization.

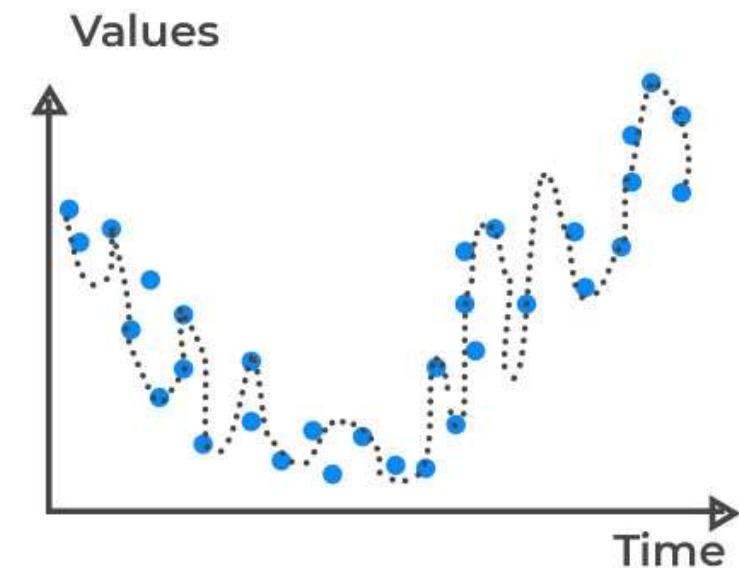
underfitting (high bias): didn't learn the data well enough



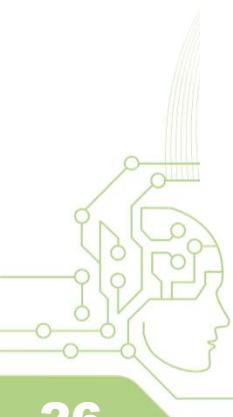
Underfitted
(High bias error)



Good Fit/R robust
(Balance between
bias and variance)



Overfitted
(High variance error)



o <https://analystprep.com/study-notes/cfa-level-/addressing-methods-method/overfitting-quantitative/>

Difference: data science, artificial intelligence, machine learning & deep learning

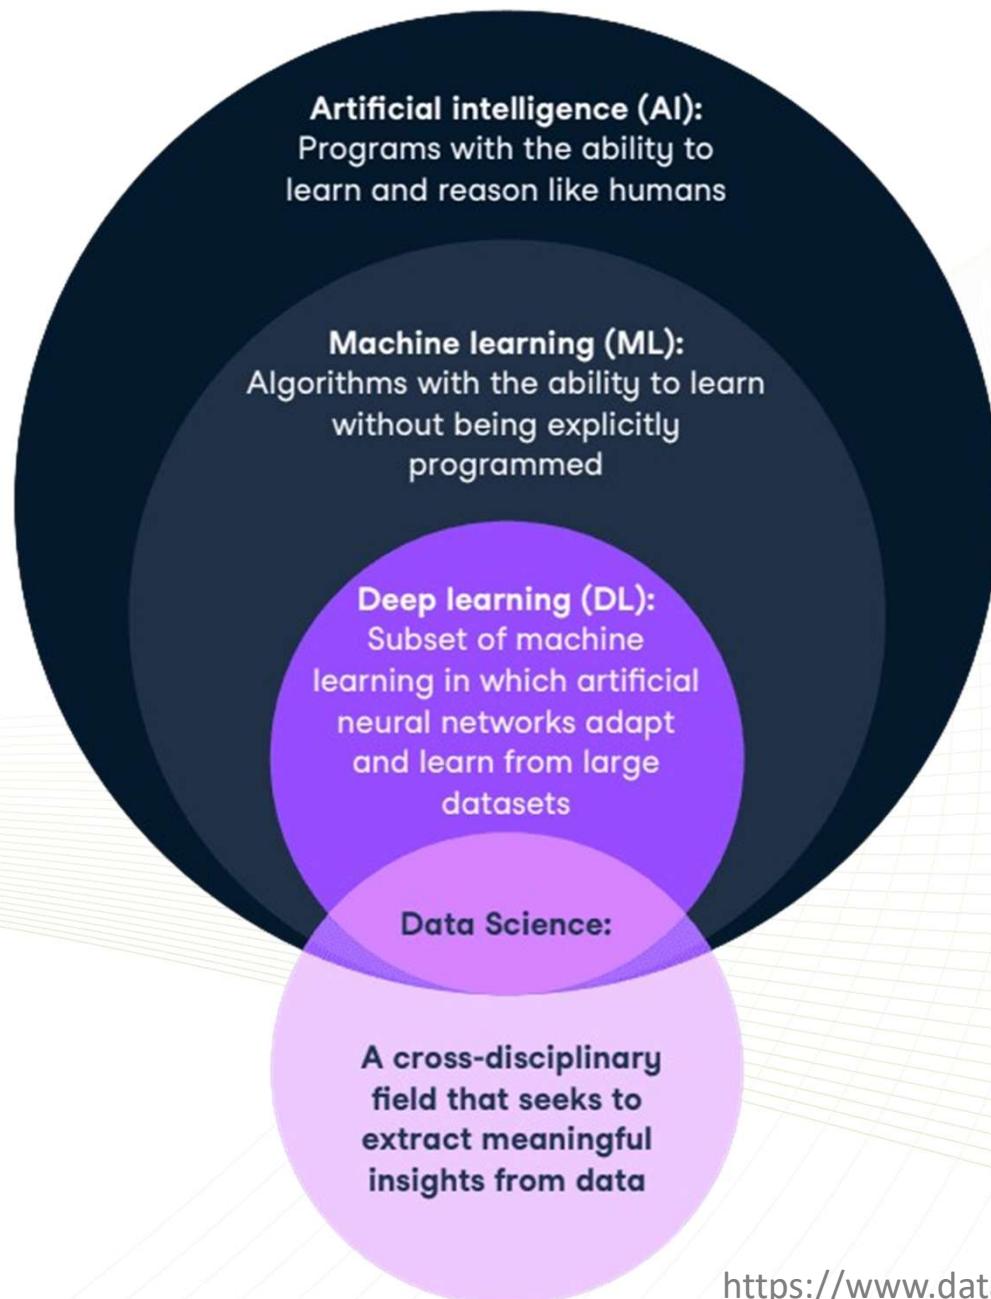
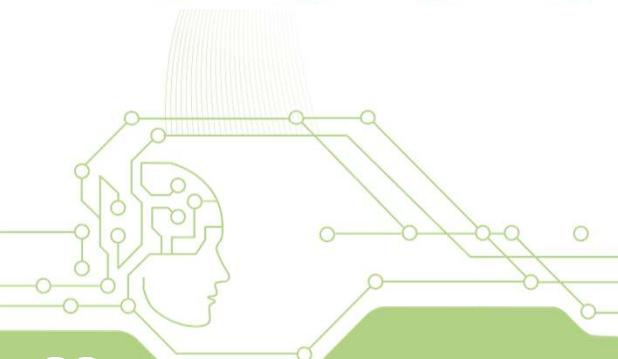
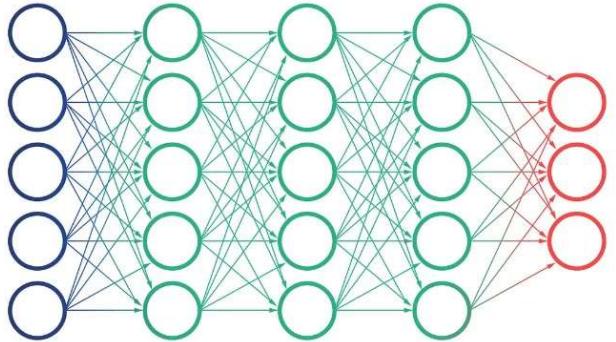
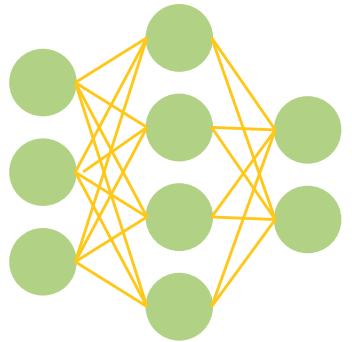
Artificial intelligence refers to computer systems that can behave **intelligently**, **reason**, and **learn** like humans.

Machine learning is a subset of artificial intelligence focused on developing algorithms with the ability to learn without explicitly being programmed.

Deep learning is a subset of machine learning. It is responsible for many of the awe-inspiring news stories about AI in the news (e.g., self-driving cars, ChatGPT). Deep learning algorithms are inspired by the brain's structure and work exceptionally well with unstructured data such as images, videos, or text.

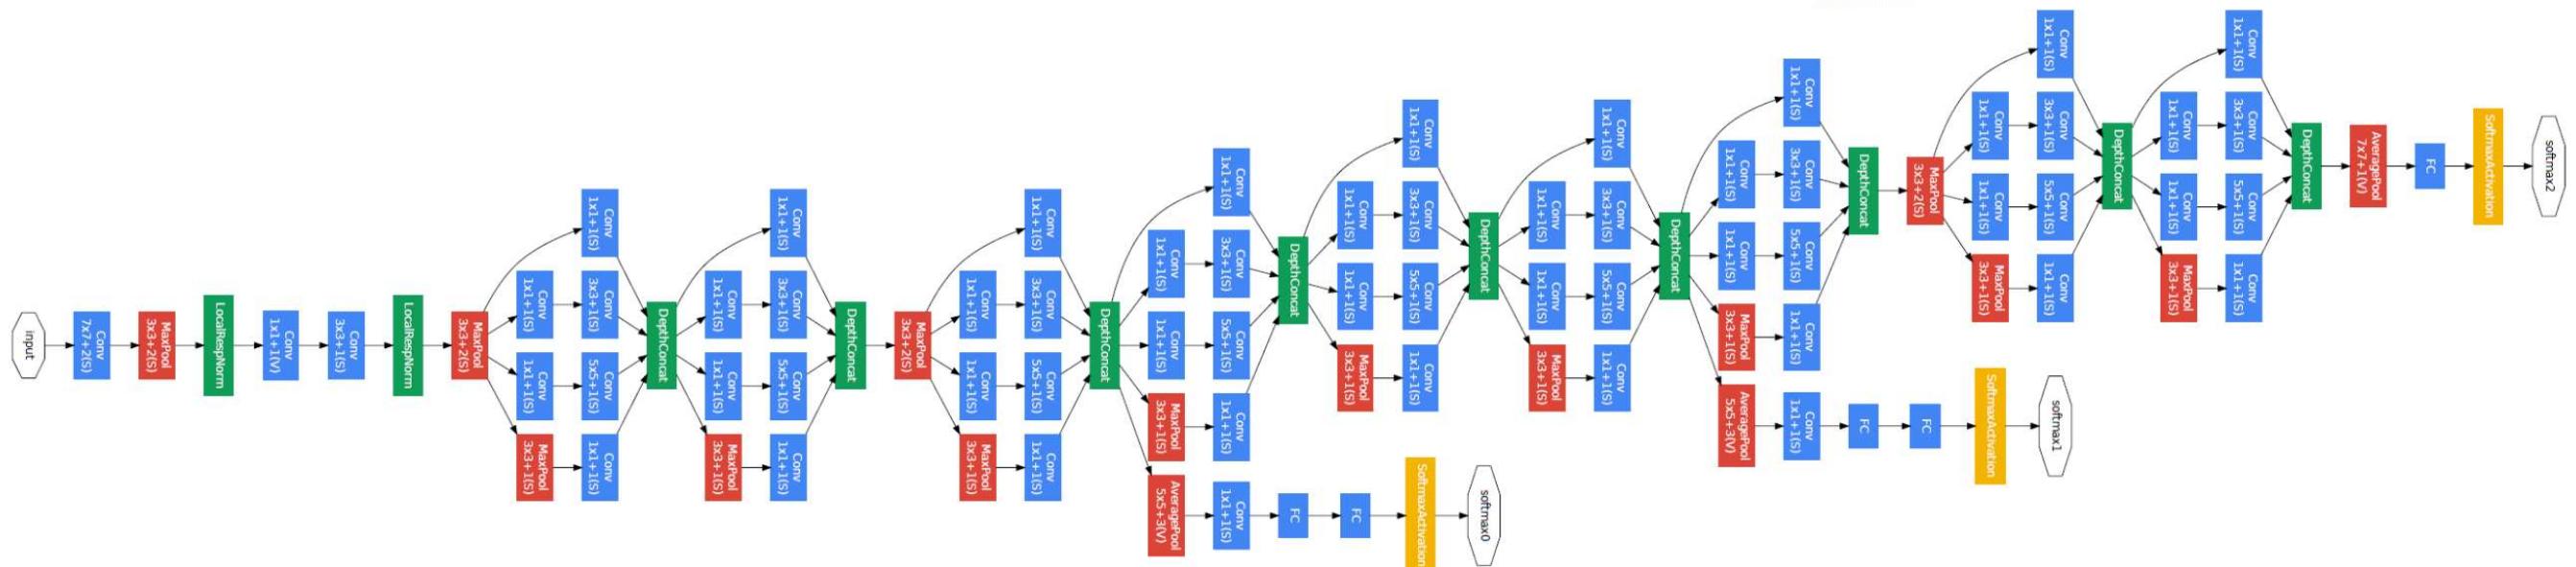
Data science is a cross-disciplinary field that uses all of the above, amongst other skills like data analysis, statistics, data visualization, and more, to get insight from data.



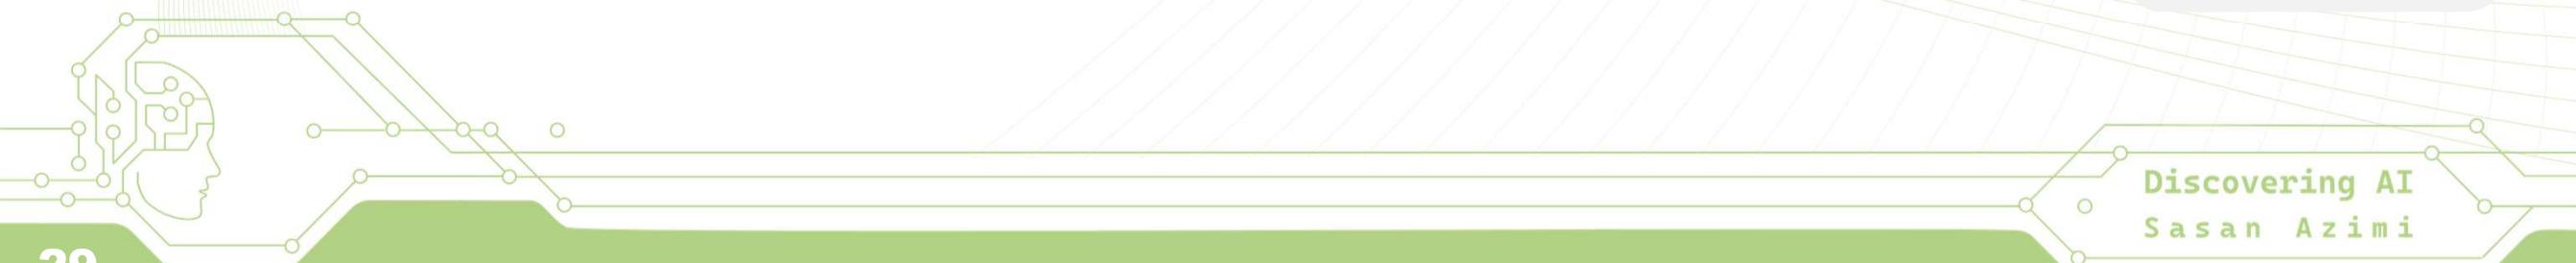


<https://www.datacamp.com>

GoogLeNet



Convolution
Pooling
softmax
Others



Some famous Neural Network Architectures

LeNet5

First convolutional neural network, leading role at the beginning of the deep learning - yann lecun - 1994

Dan Ciresan Net

First implementation of GPU neural -2010. 9 layers - NVIDIA GTX 280

AlexNet

Won the challenging competition of imagenet

Overfeat

Derivative of AlexNet

VGG

Used smaller 3×3 filters in each convolutional layers – based on LeNet by Oxford

GoogLeNet

first architecture to decrease the burden of computation of deep neural networks.

ResNet

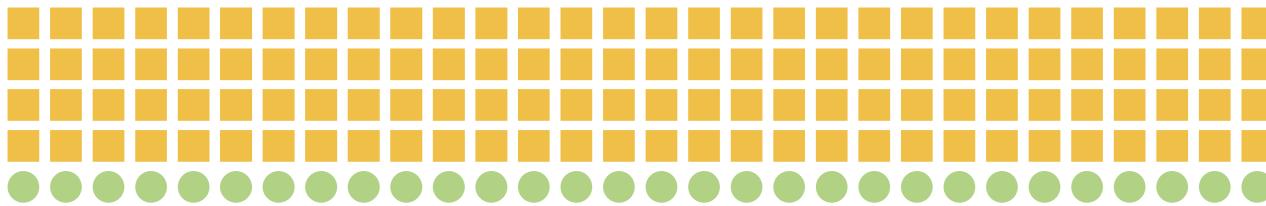
Straightforward, More than a hundred and thousand layers trained first time

SqueezeNet

Inception and ResNet's concepts have been re-hashed in SqueezeNet

ENet

Very light-weight, low computations and parameters. Combining modern architectures' features.



train

validation

test

x_train
y_train

x_test
y_test

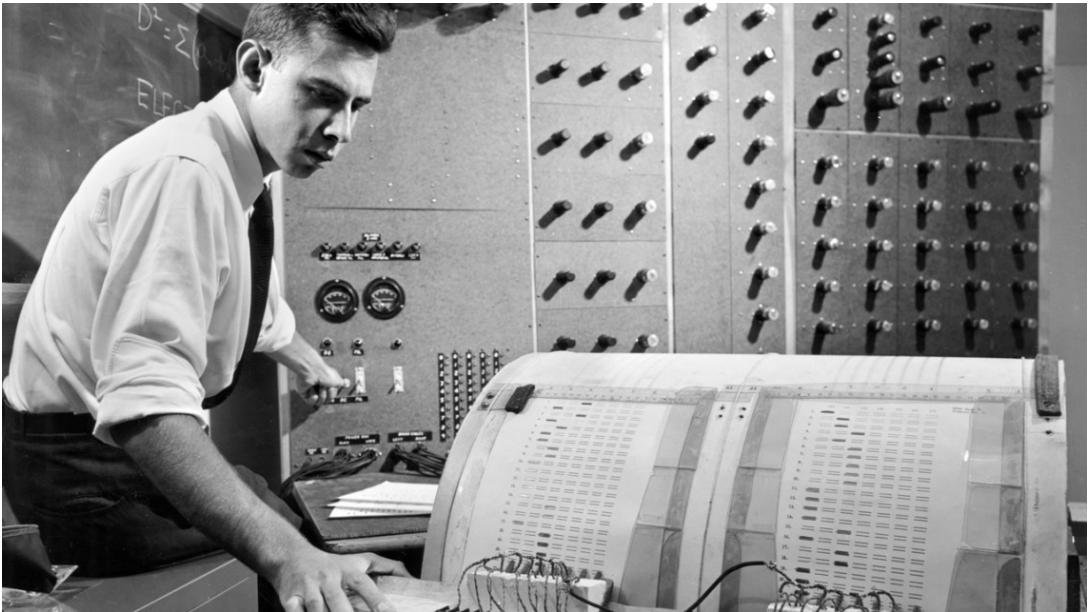
Train

Email	Label
Click here to claim your prize!	Spam
What's new?	Not spam
Hang out later?	Not spam
...	...

Test

Email	Label
Pick up groceries	Not spam
Free free free!	Spam
Android 5.0 questions	Not spam
...	...





Frank Rosenblatt
psychologist



First Model of **Perceptron** for Supervised learning

Obtaining the weights using the input data by the neural network itself

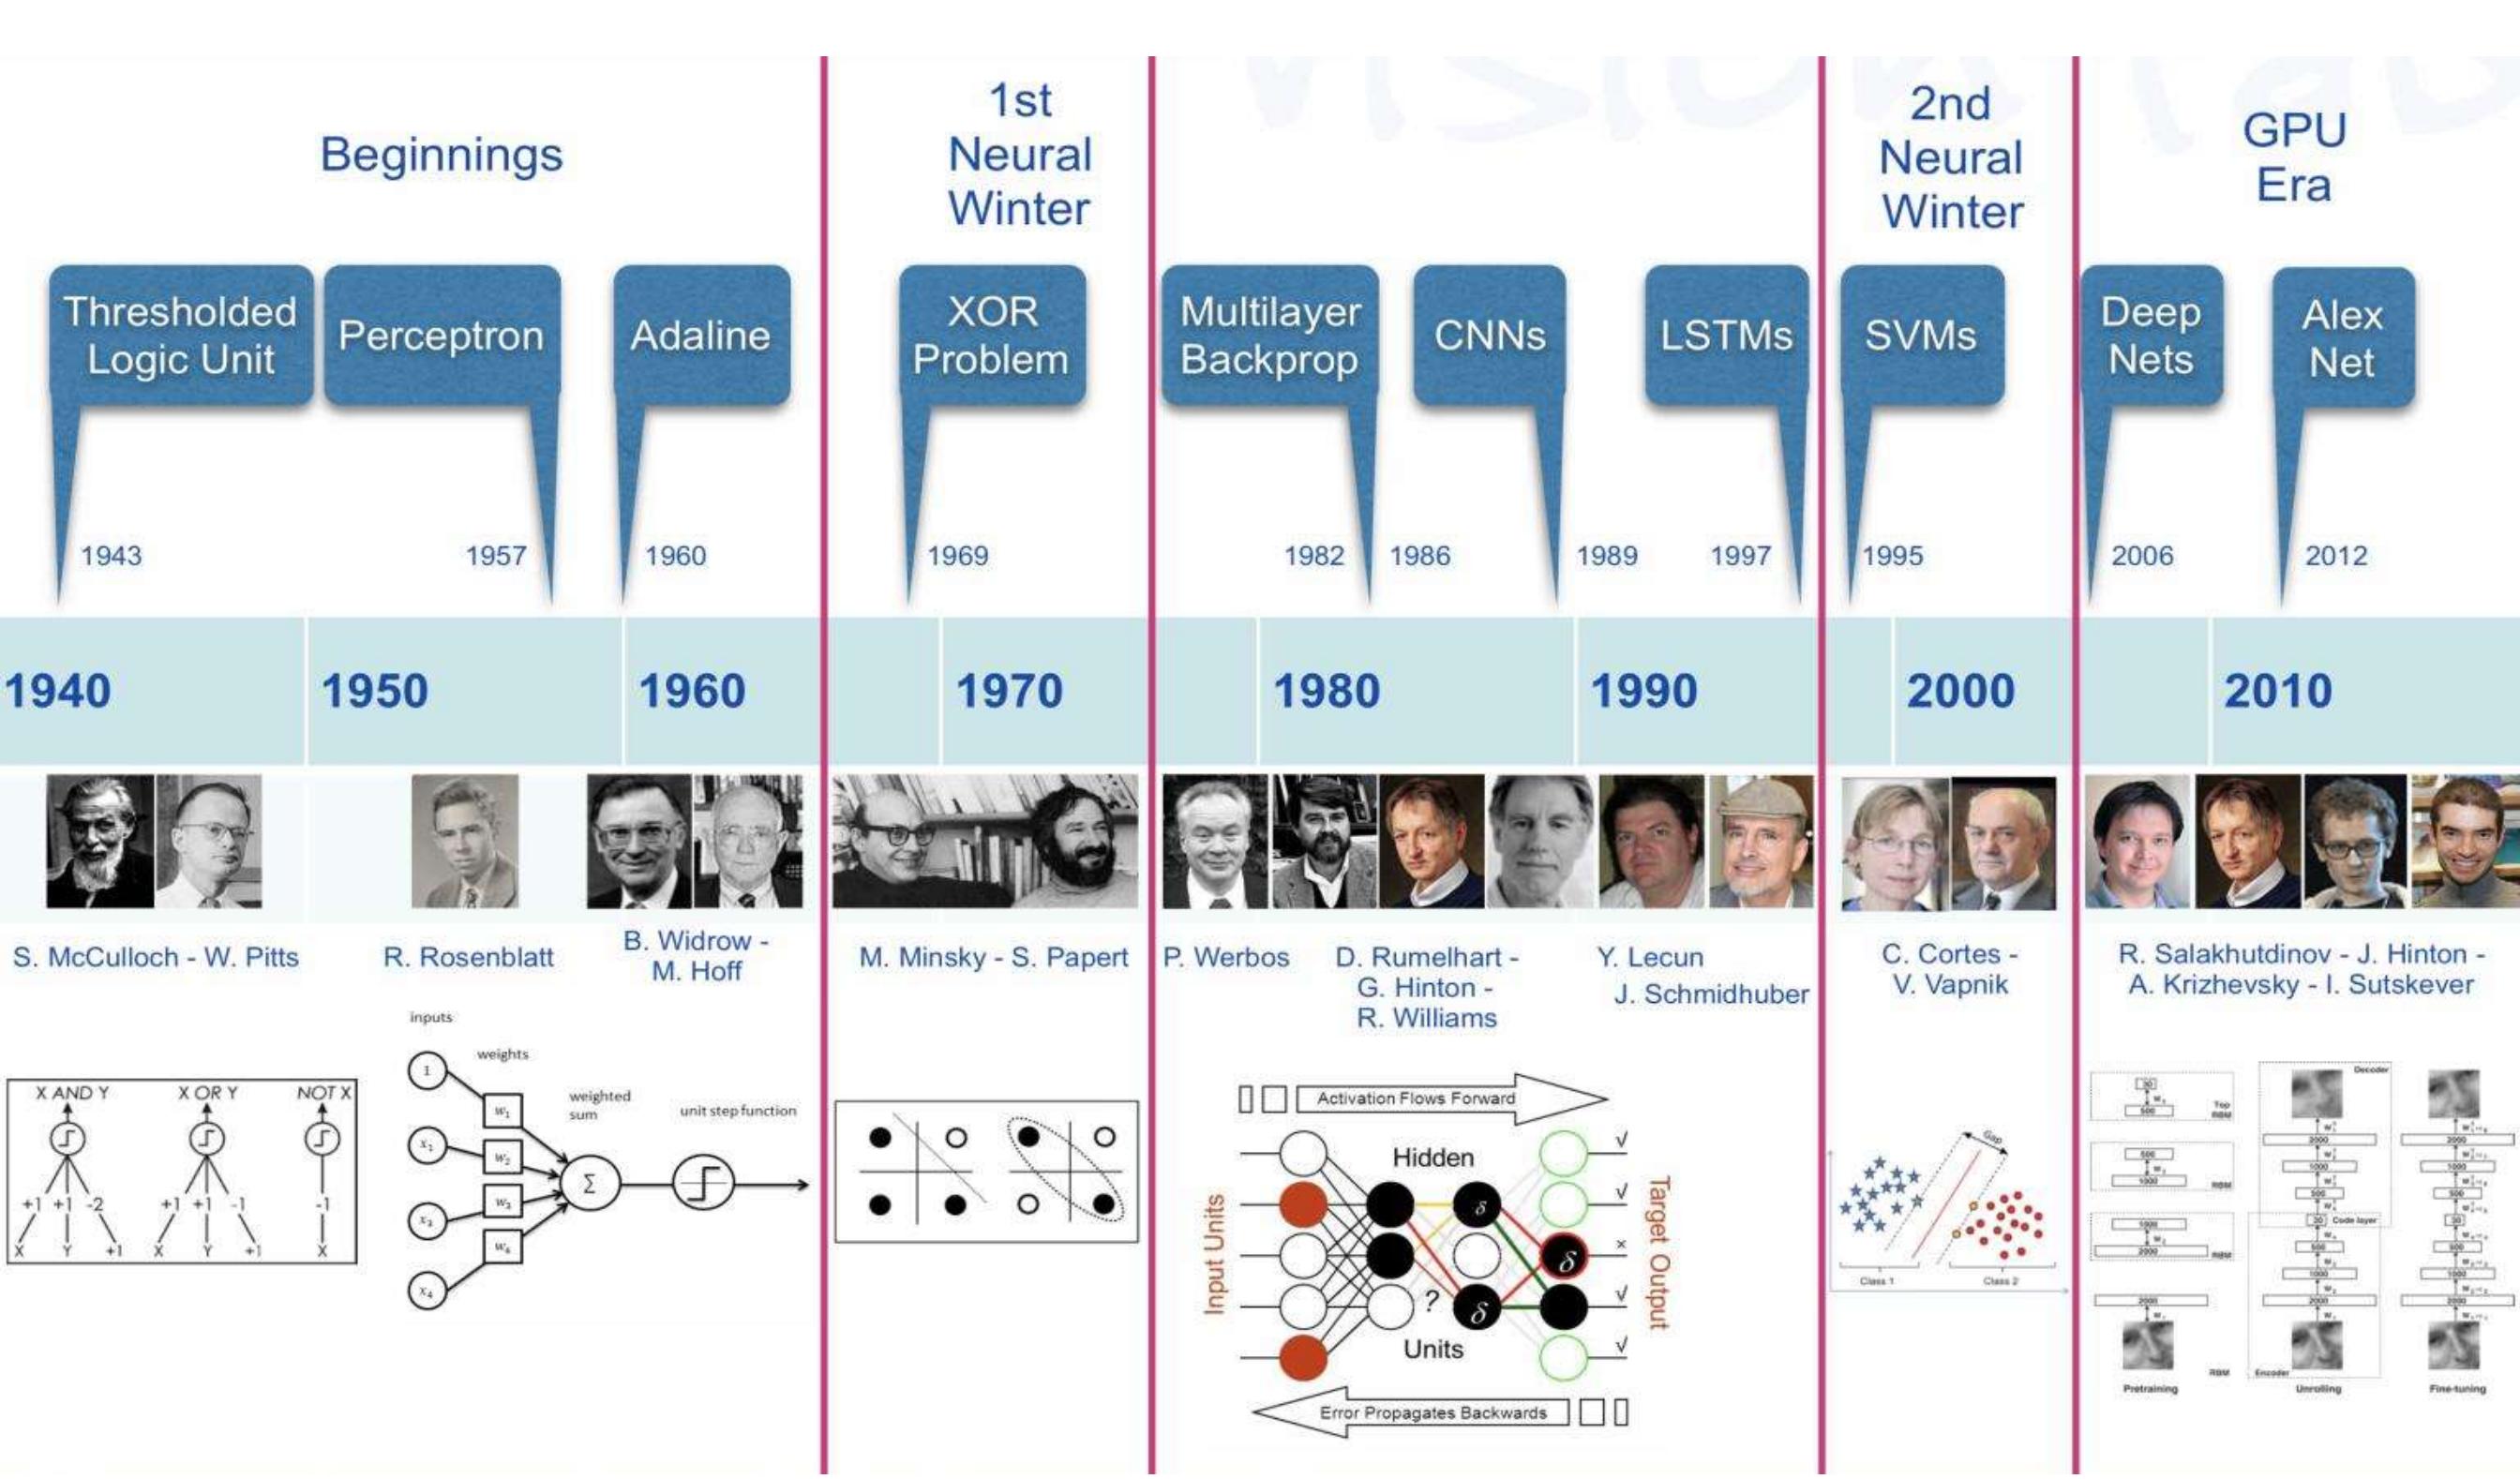
Similar to the todays perceptron networks

IEEE annual award

After 1952

- Early enthusiasm, great expectations (1952–1969)
- Expert systems (1969–1986)
- The return of neural networks (1986–present)
- Probabilistic reasoning and machine learning (1987–present)
- Big data (2001–present)
- Deep learning (2011–present)

Reference: www.simplypsychology.org



Implementing Machine Learning

Examine pillars of a practicing AI team

Bettering Machine Learning Model Management

State of tools – understanding machine learning stacks

Machine Learning Methods and Algorithms

Developing Validation Sets

Developing Training Sets

Accelerating Training

Encoding Domain Expertise in Machine Learning

Automating Data Science

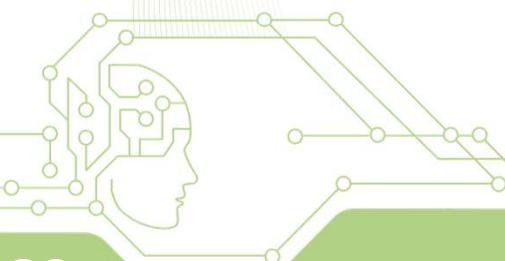
Deep Learning

- 1. Business case
 - 2. Domain expertise
 - 3. Data science
 - 4. Algorithms
 - 5. Application integration
-
- 1. Decision Trees
 - 2. Support Vector Machines
 - 3. Regression
 - 4. Naïve Bayes Classification
 - 5. Hidden Markov Models
 - 6. Random Forest
 - 7. Recurrent Neural Networks
 - 8. Convolutional Neural Networks
 - 9. ...



Types of Algorithms used in Deep Learning

1. Convolutional Neural Networks (**CNNs**)
2. Long Short Term Memory Networks (**LSTMs**)
3. Recurrent Neural Networks (**RNNs**)
4. Generative Adversarial Networks (**GANs**)
5. Radial Basis Function Networks (**RBFNs**)
6. Multilayer Perceptrons (**MLPs**)
7. Self Organizing Maps (**SOMs**)
8. Deep Belief Networks (**DBNs**)
9. Restricted Boltzmann Machines (**RBMs**)
10. Autoencoders



Attention mechanisms

Neural network focus on the most relevant parts of the input or output for a given task.

Different types

Compute a score or a weight for each element in the input or output sequence, and then use these scores to aggregate the elements into a context vector.

The attention mechanism borrows from the signal processing mechanism of the human brain.

Its application in convolutional neural networks is reflected in generating differences in the importance of each piece of input information



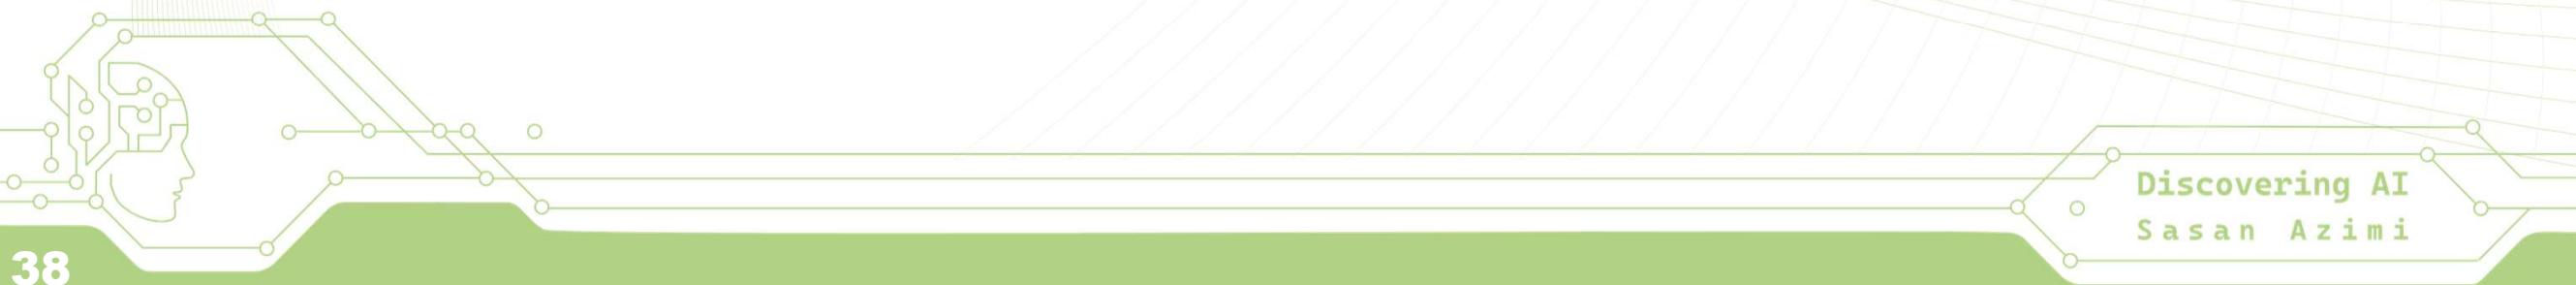
Attention mechanisms

Examples of attention mechanisms:

Scaled Dot-Product Attention: This is a type of attention that computes the similarity between a query vector and a set of key vectors, and then uses these similarities to weight the corresponding value vectors. This is used in the Transformer model

Additive Attention: Computes the score for each element by concatenating the query vector and the key vector, and then passing them through a feed-forward network. This is also known as Bahdanau attention.

Channel Attention: Computes the importance of each channel in a feature map, and then uses these weights to rescale the feature map. This is used in Squeeze-and-Excitation Networks.



Ultra realistic alien artic fractal fungal plant tree
proportinal geometry dynamic composition,
asymmetrical composition,golden ratio,
photorealism, Lens with Nikon d850, 28mm, Film
Light, Overall Light, Ethereal Light, Depth of Field,
f/2.8, Ultra HD, 128k, 3D Shadows, Tone
Mapping, Ray Traced Global Illumination, Super
Resolution, Gigapixel, Color Correction,
Retouching, Enhancement, blue sky, PBR,
Blender, V-ray, Procreate, Unreal Engine 5,
Cinema 4D, ROMM RGB, Adobe After Effects,
3DCG, VFX, SFX, FXAA , award-winning
photography, natural lighting, to p light, studio
light, rim light, Volumetric Lights , RTX Renderer,
High Qual enviornment, Extreme Realism, RTX
On, Realistic, Hyper Realism, Realistic Textures,
Amazing Props, Making Amazing Visual Effects,
Crazy Details, High Details, Intrinsic Details, Post
Production, cool palette, depth of field, IMAX,
ALEXA LF, --ar 3:2 --s 999 --w 2311



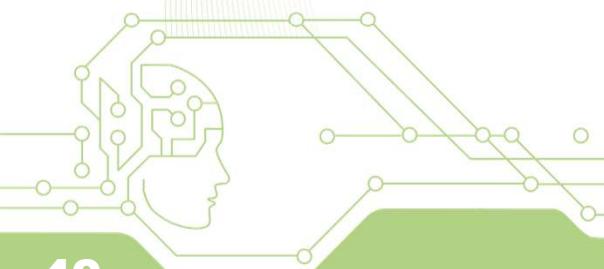
Understanding Generative Adversarial Networks (GANs)



Ilya Sutskever
OpenAI

**Self-supervised
learning**

Has taken the world



Machine learning sucks!



The most interesting idea
proposed in the field of artificial
intelligence in the last ten years

Yann LeCun – Meta-NYU



Ion
Goodfellow

